


# Hardware Accelerated Proofs

Erdinc Ozturk, Justin Drake, Simon Peffers, Sean Gulley, Kelly Olson

05.14.20

# Outline

- Why acceleration
- RSA/ECC primitives
- Algorithms
- Platforms - CPU, GPU, FPGA, ASIC
- Performance comparisons
- SNARKs



Intro

# Why Accelerate?

Conventional cryptography is well-studied

Innovation in cryptography is enabling new capabilities

- VDFs, proofs of exponentiation, SNARKs, accumulators, ...

- Unbiasable randomness, privacy, scalability, compression (compute/storage), ...

Most of these are computationally intensive

Acceleration enables:

- Higher security: larger key size = increased security

- Greater scale

- Improved user experience

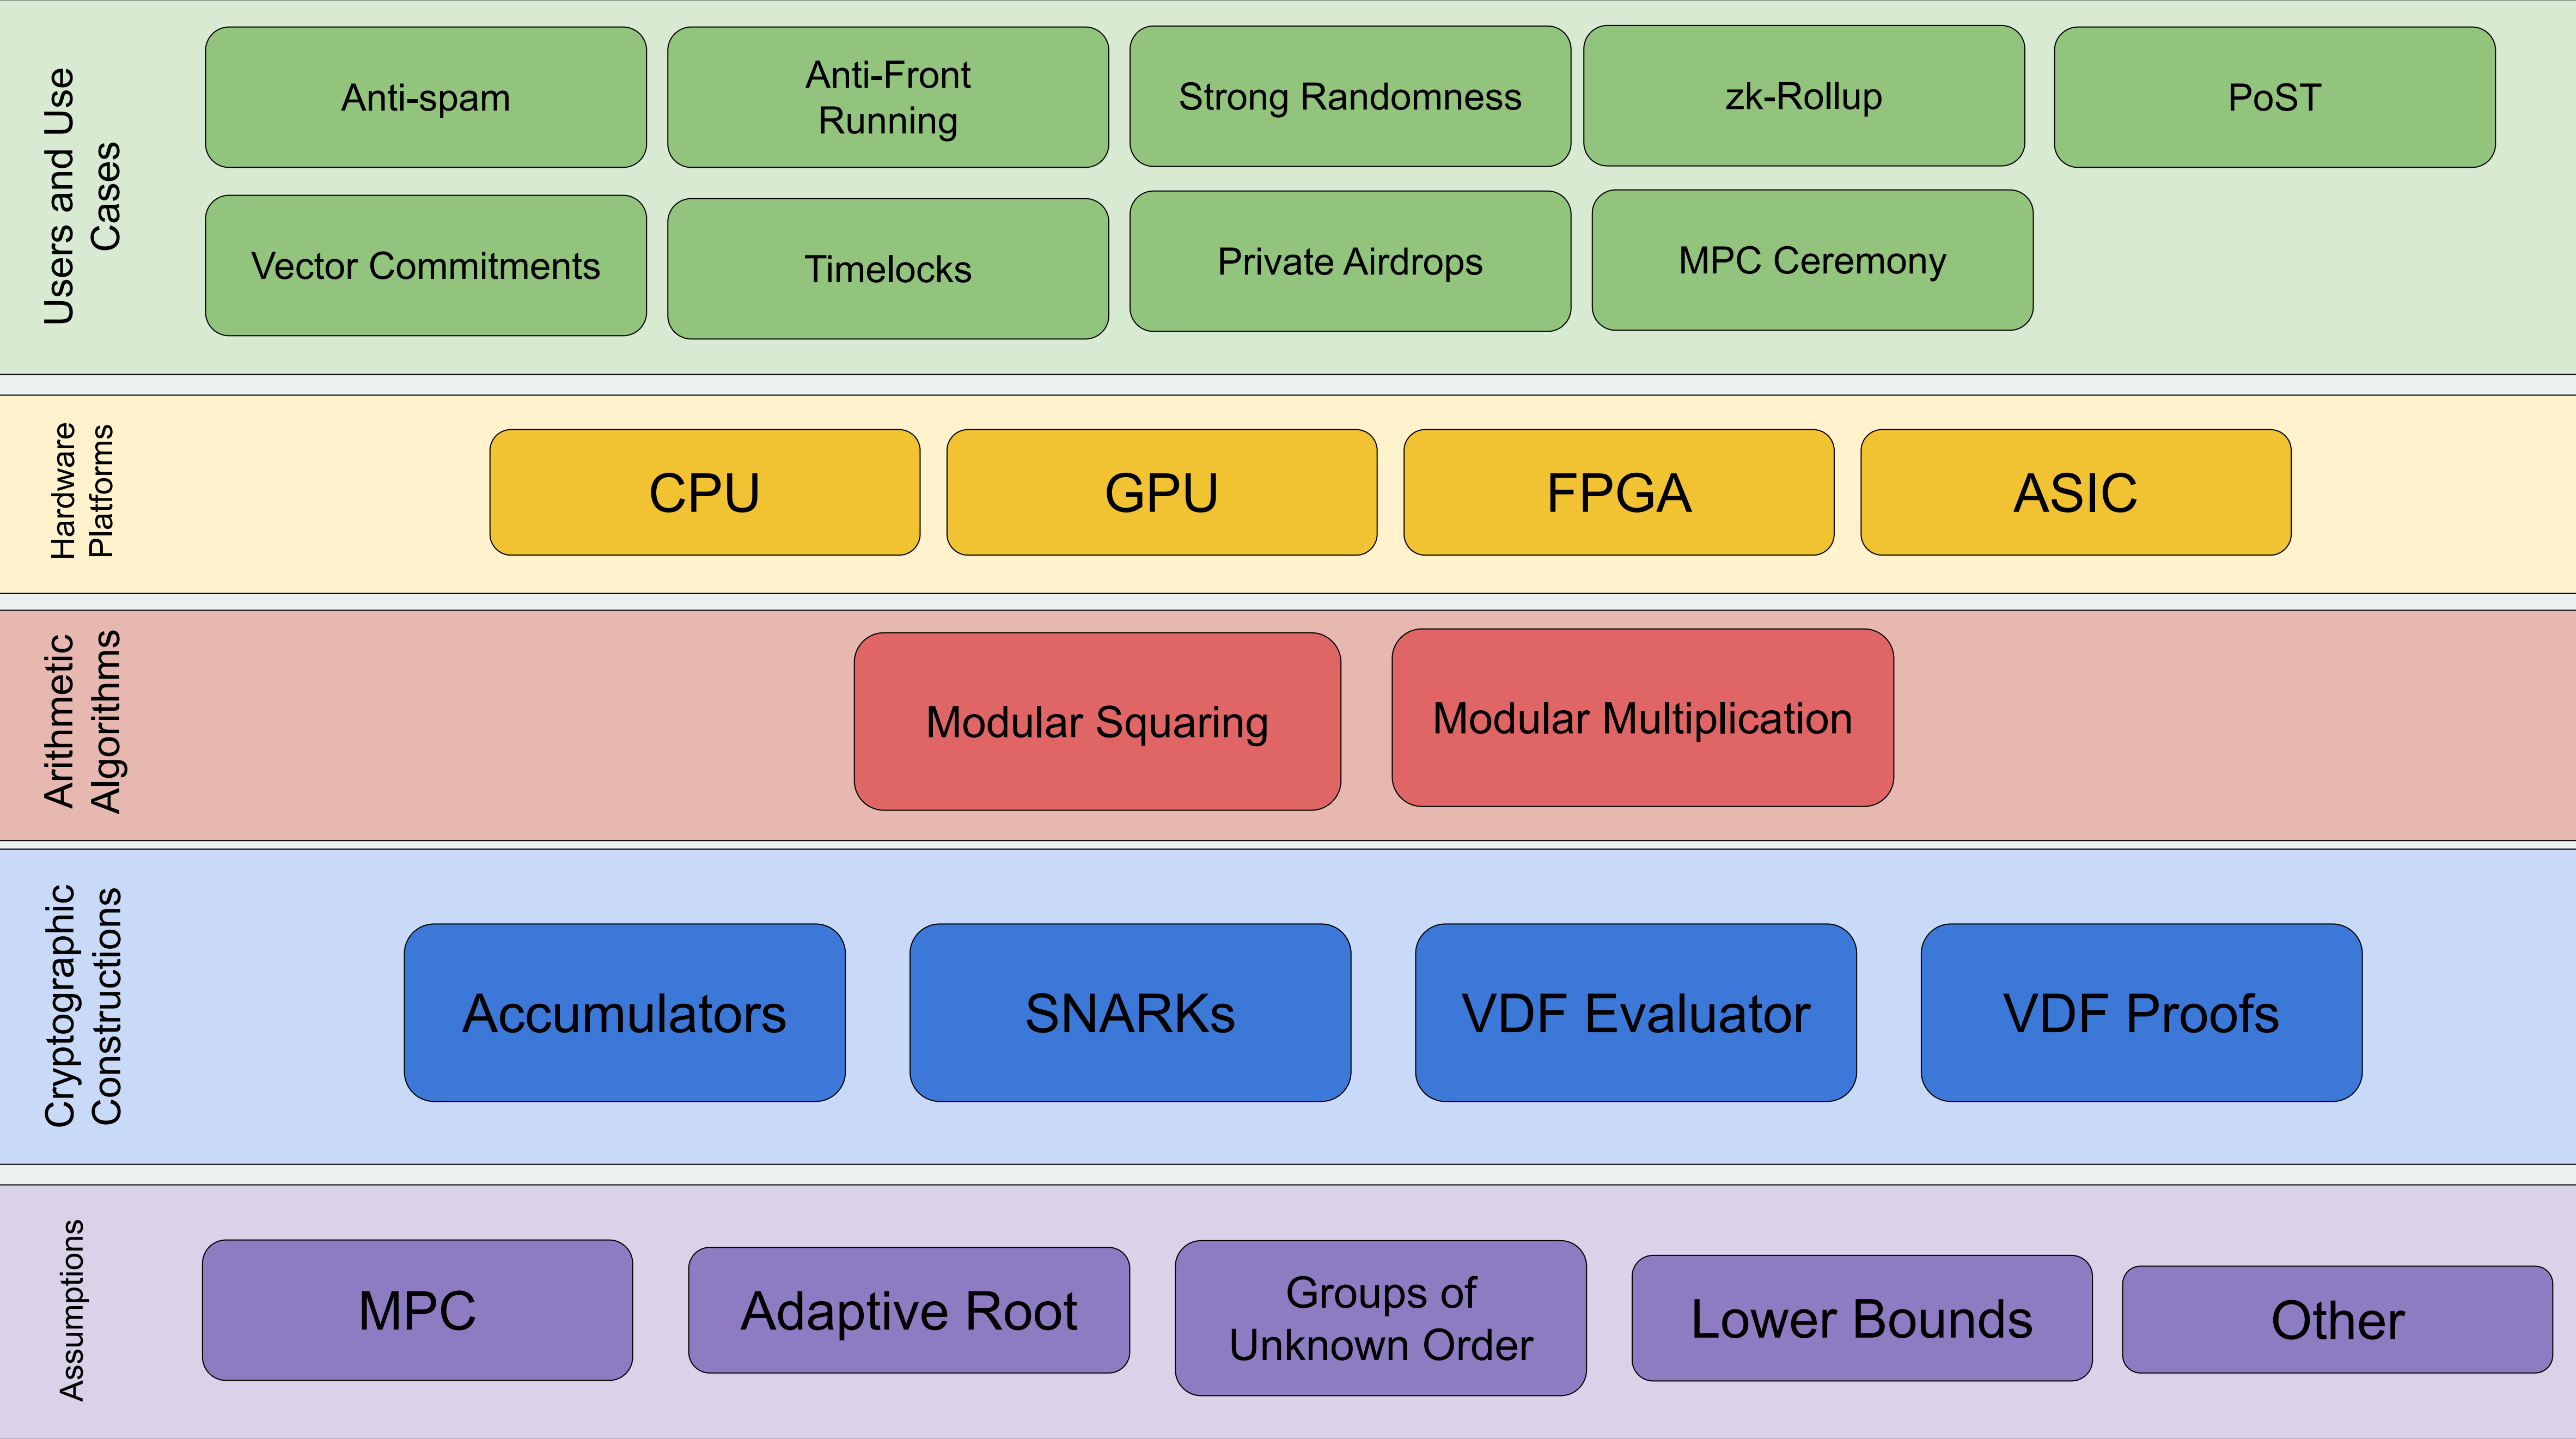
- Lower costs

The same large integer arithmetic is fundamental across cryptography - RSA, ECC, etc.

# Process

# **Recipe for Acceleration**

- Define the use case
- Benchmark / profile to find bottlenecks
- Identify the underlying key operations to accelerate
- Select target platform(s) based on required performance characteristics
- Carefully map the operations onto the capabilities provided by the target platform



# RSA Primitives

RSA consists of modular exponentiation

Exponentiation:  $x^y$       <- square and multiply

VDF:  $x^{(2^t)}$       <- repeated square

Many different exponentiation techniques

Example: left-to-right square and multiply

Example: left-to-right square and multiply with windowing -> precomputed lookup tables

All rely on these large integer modular arithmetic primitives:

Square

Multiply

Add

Subtract



# ECC Primitives

ECC consists of scalar point multiplication

$$Q = k * P$$

All exponentiation techniques can be used for scalar point multiplication.

High level primitives are replaced:

Modular squaring             $\rightarrow$  point doubling

Modular multiplication    $\rightarrow$  point addition

Point addition and Point doubling operations rely on same large integer modular arithmetic primitives:

Smaller operands

One major caveat: point at infinity

# Latency vs. Throughput Oriented Problems

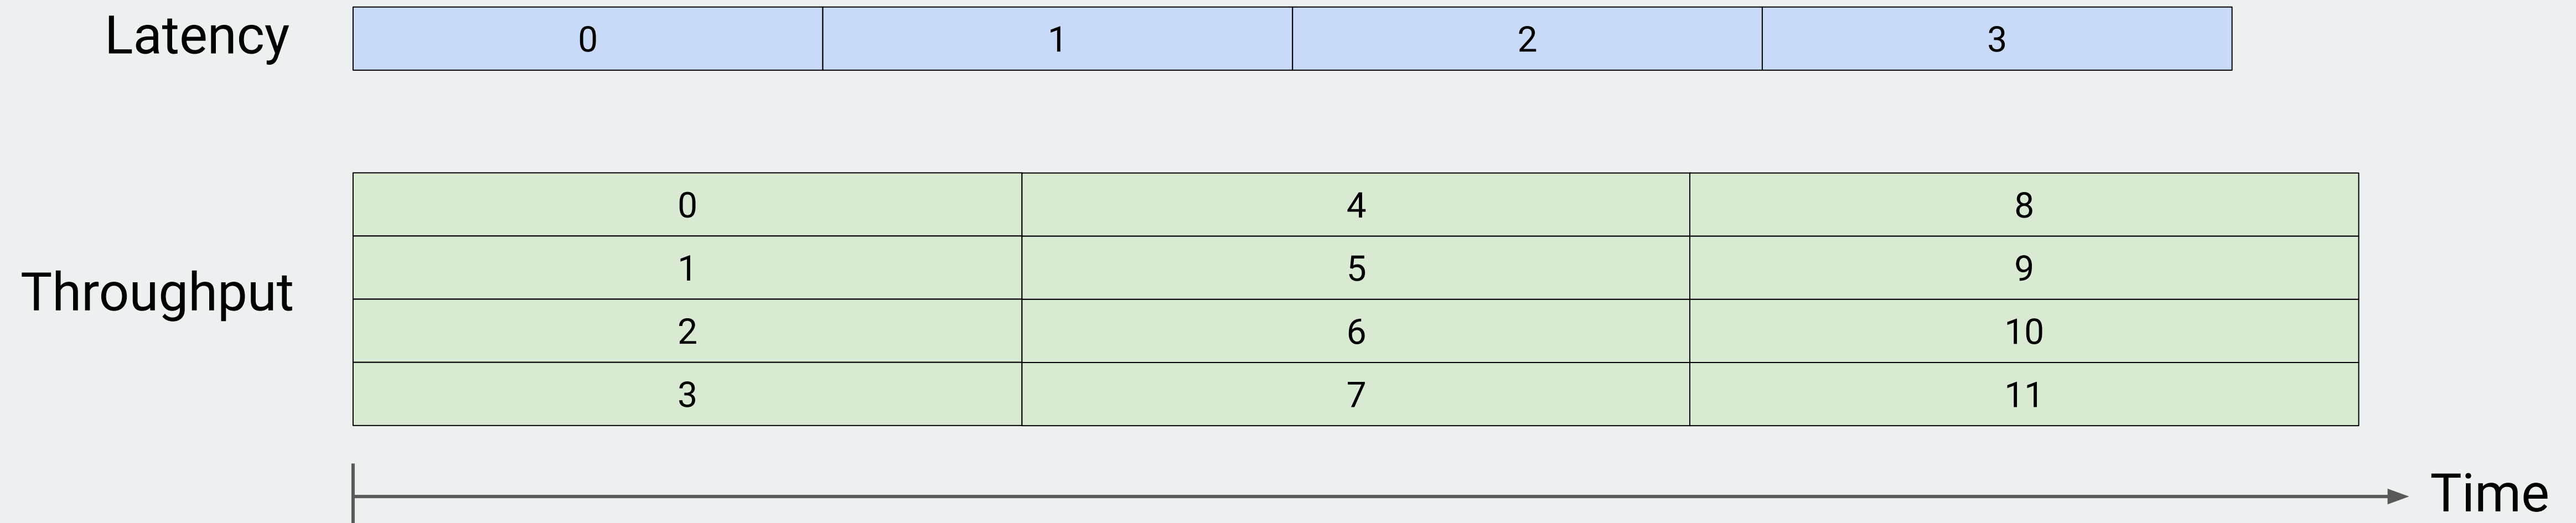
Latency oriented is performing an individual task as quickly as possible

Example - VDF Evaluation

Throughput oriented trades off time of an individual task for performing multiple tasks in parallel

Example - VDF Proof

Choose the appropriate algorithm and platform based on problem priority



# Background

# Performance Impacts - Algorithms

Implementation algorithms can lead to orders of magnitude difference in performance

## Examples

Modular Squaring/Multiplication

Schoolbook, Comba, Karatsuba, Schonhage-Strassen

Montgomery, Barrett

Modular Exponentiation

Binary square and multiply, Montgomery Ladder, Sliding window, wNAF, Addition Chain

Modular Multi-Exponentiation

Pippenger, Bos-Coster, Lim-Lee, Comb, Shamir

Exponentiation/Multi-Exponentiation Considerations

Fixed base, fixed exponent

# Multi-exponentiation

Many different flavors:

Different Base, Different Exponent:  $g_1^{x_1}$   $g_2^{x_2}$  ...  $g_n^{x_n}$

Same Base, Different Exponent:  $g^{x_1}$   $g^{x_2}$  ...  $g^{x_n}$

Different Base, Same Exponent:  $g_1^x$   $g_2^x$  ...  $g_n^x$

ECC:

Different Base, Different Exponent:  $k_1 * P_1$   $k_2 * P_2$  ...  $k_n * P_n$

Same Base, Different Exponent:  $k_1 * P$   $k_2 * P$  ...  $k_n * P$

Different Base, Same Exponent:  $k * P_1$   $k * P_2$  ...  $k * P_n$

# Multi-exponentiation

Example:

Calculate  $(k_1 * P_1) + (k_2 * P_2) + (k_3 * P_3) + (k_4 * P_4)$

	scalar index							
	7	6	5	4	3	2	1	0
k1	1	0	1	1	0	1	0	1
k2	1	1	0	1	1	0	1	1
k3	0	1	1	0	1	1	0	1
k4	1	0	1	1	0	1	1	0

Complexity for 4 bases:

Single exponentiation:    Point Doublings:  $4*7$     Point Additions: 20

Multi exponentiation:    Point Doublings: 7    Point Additions: 20

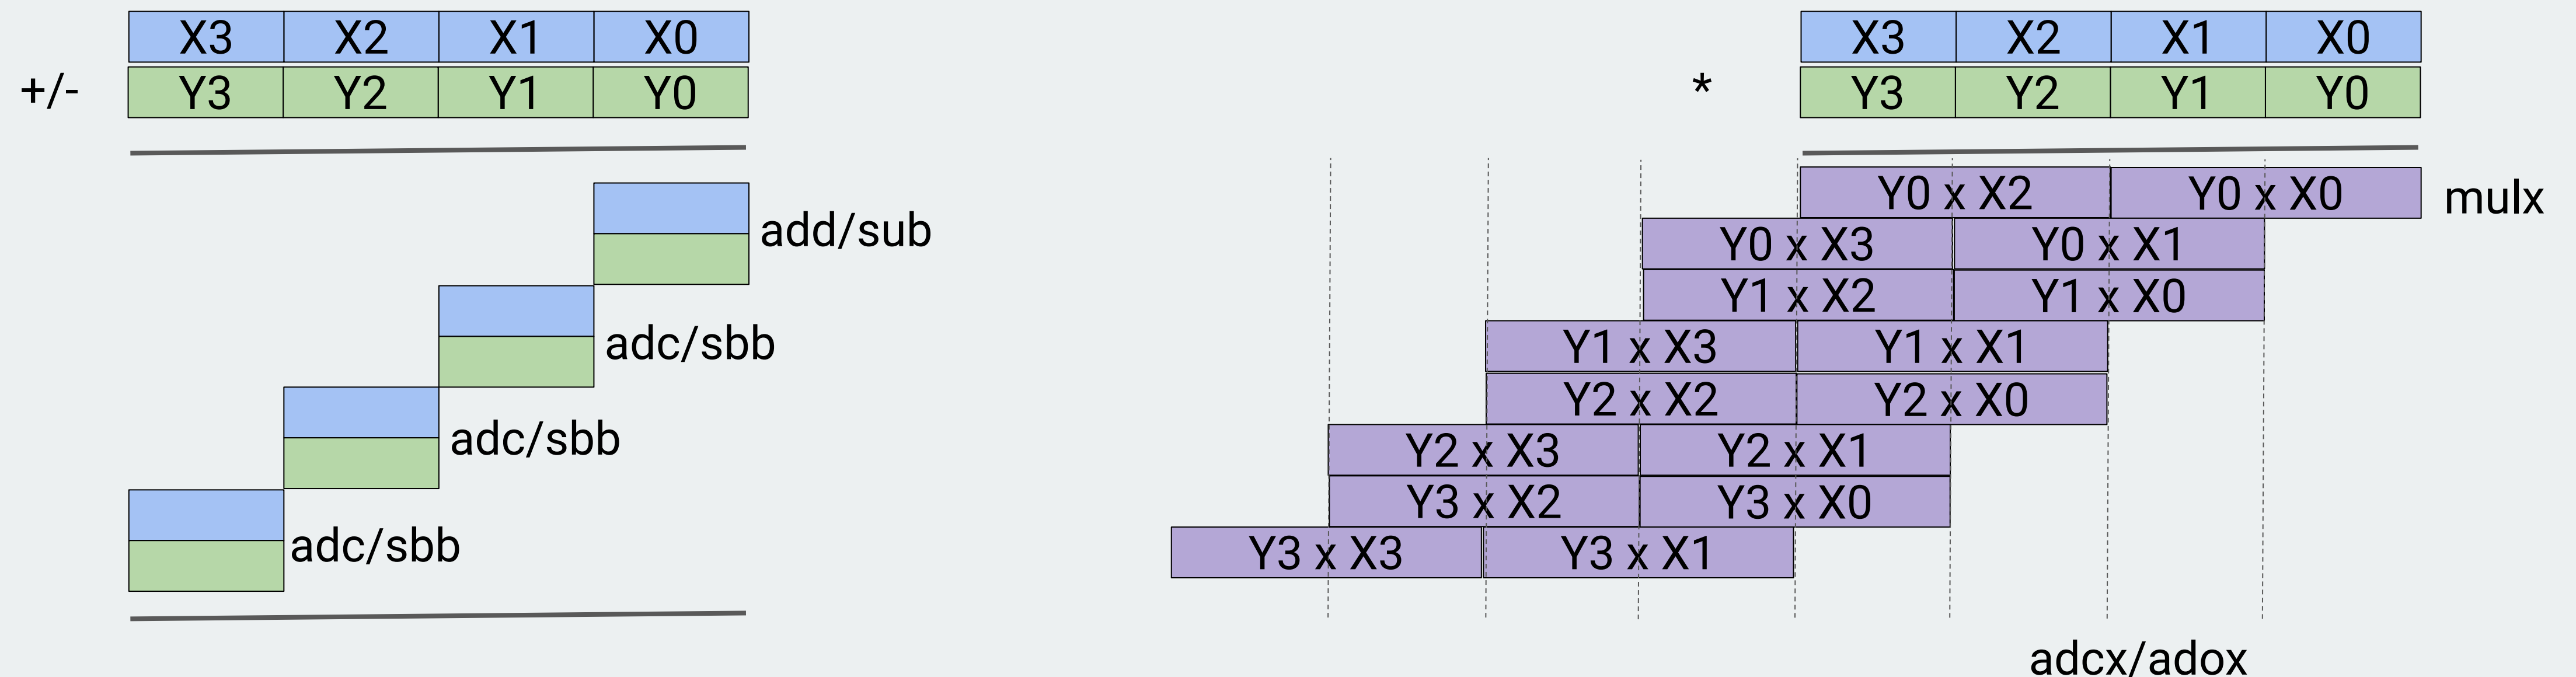
No constant time!

# Large Integer Arithmetic

RSA integer sizes typically exceed compute element word size (e.g. 8, 16, 32, 64 bits)

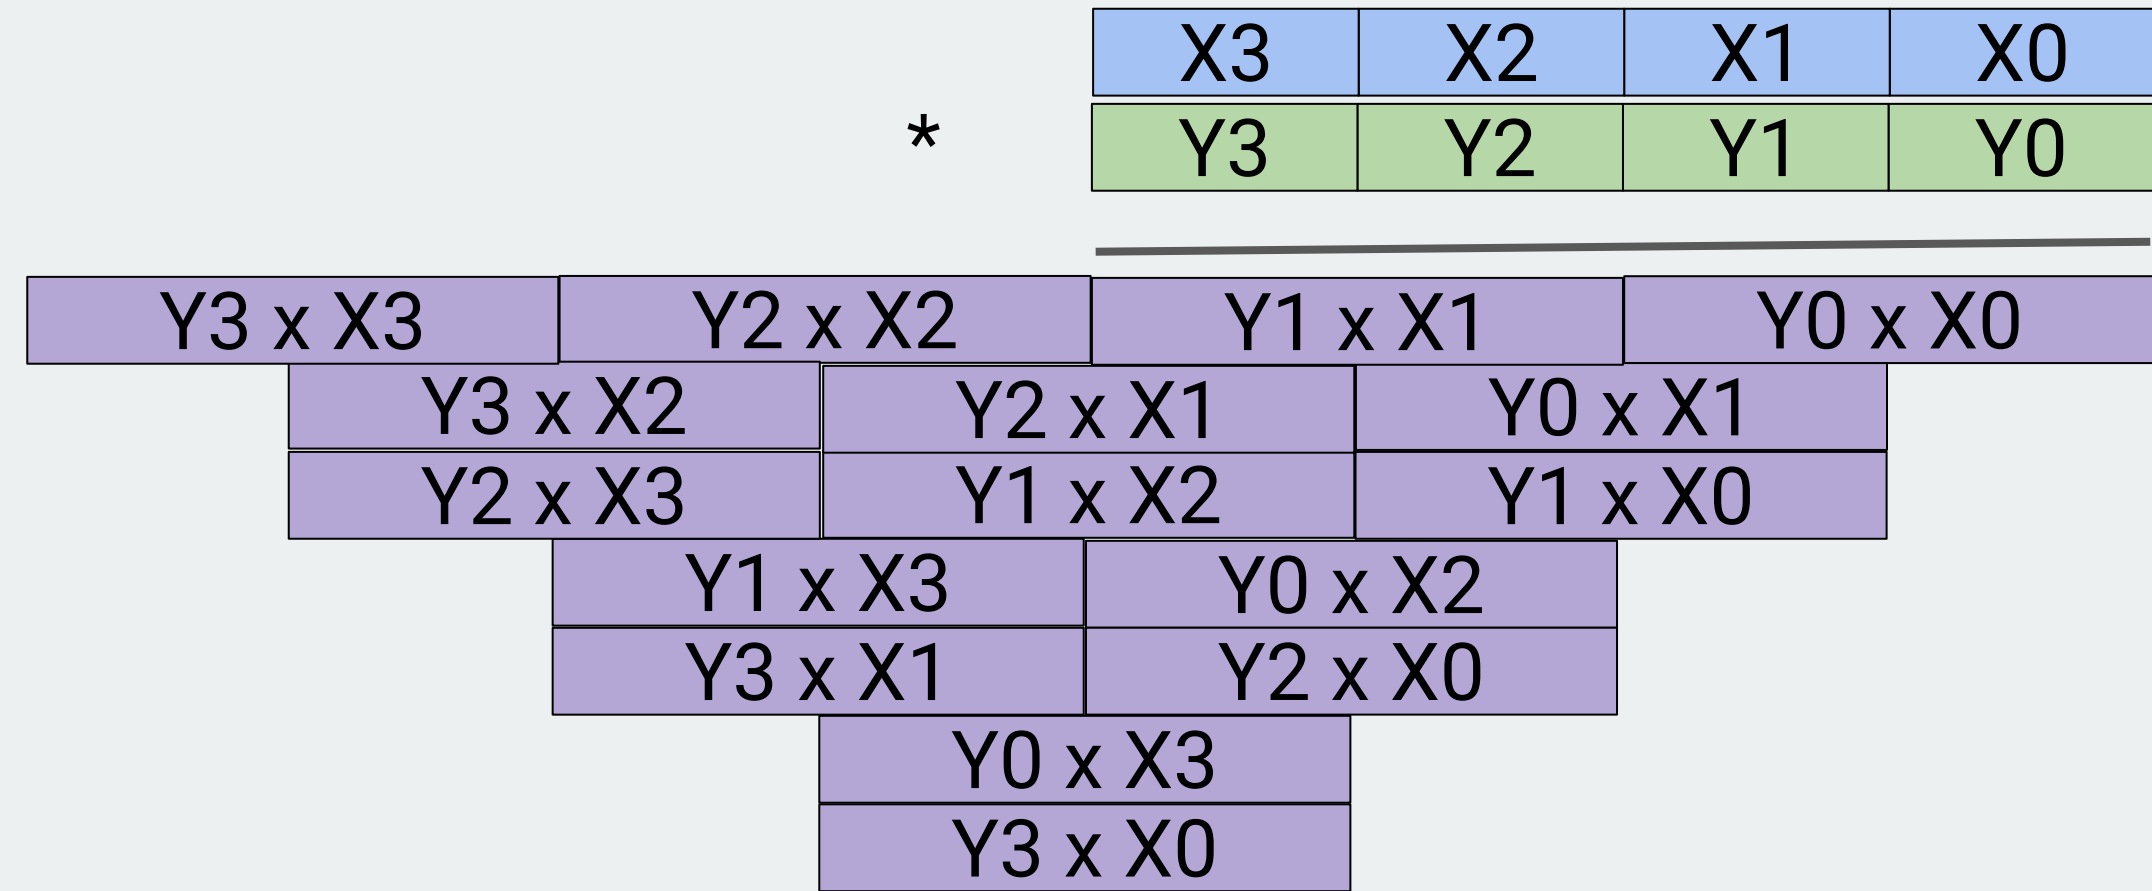
Break integer down into n limbs:  $X = (X[n-1], \dots, X[1], X[0])$

$n = \text{ceil}(\text{len}/w)$  where len is the integers bit length and w is compute word size

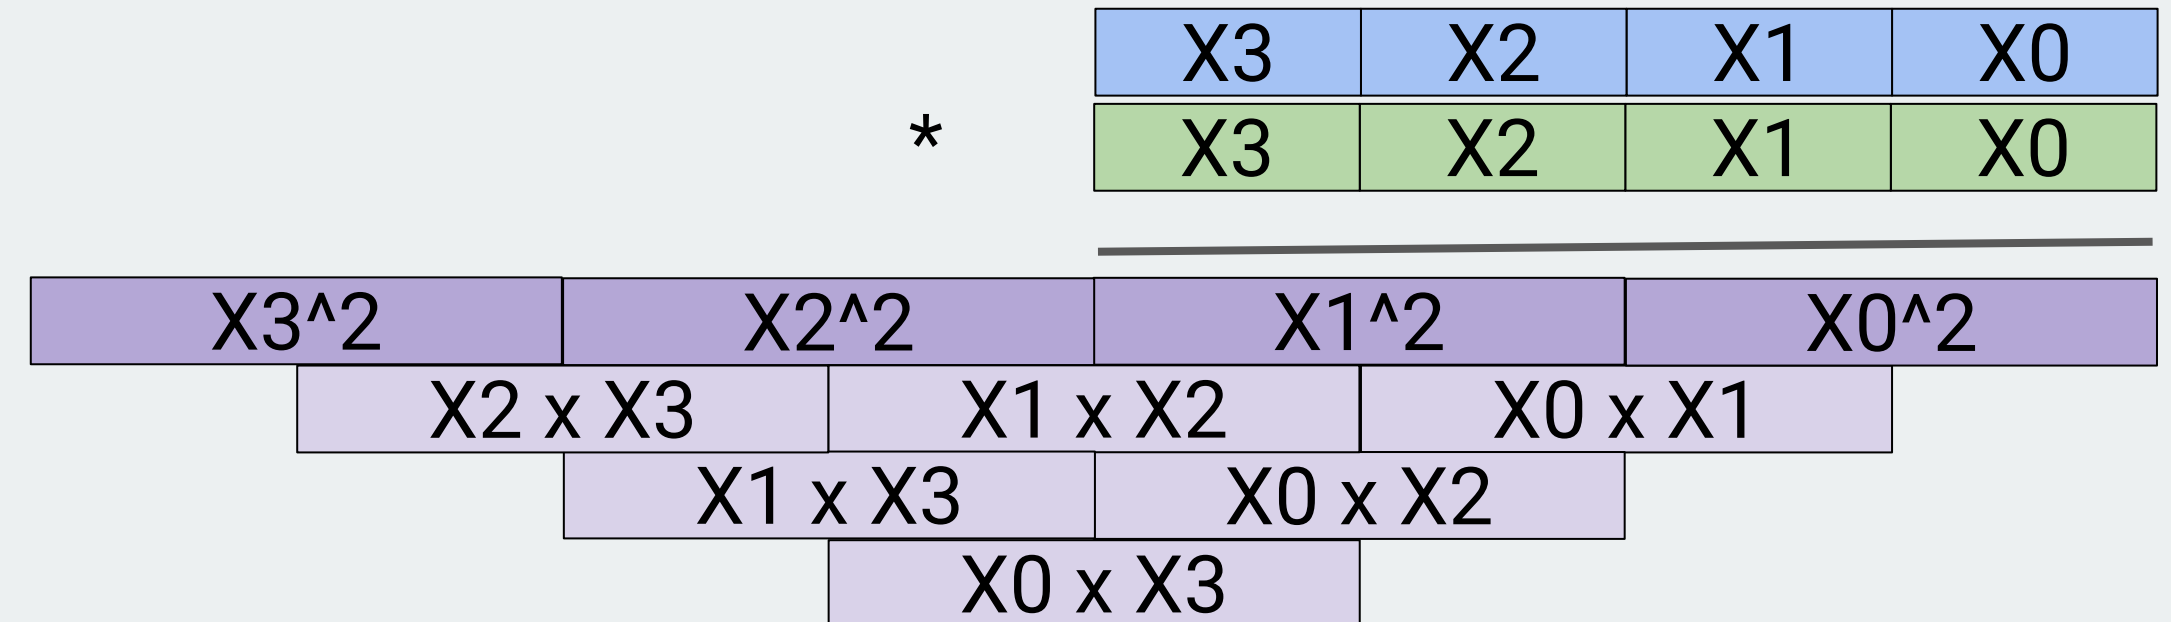


Parallel carry chain, use adox for low word accumulation and adcx for high word (or vice versa)

# Large Integer Arithmetic




Schoolbook multiply requires  $n^2$  operations




Square requires  $(n^2 + n)/2$  operations  
Light purple partial products are doubled



# Large Integer Arithmetic



Schoolbook multiply requires 1024 operations

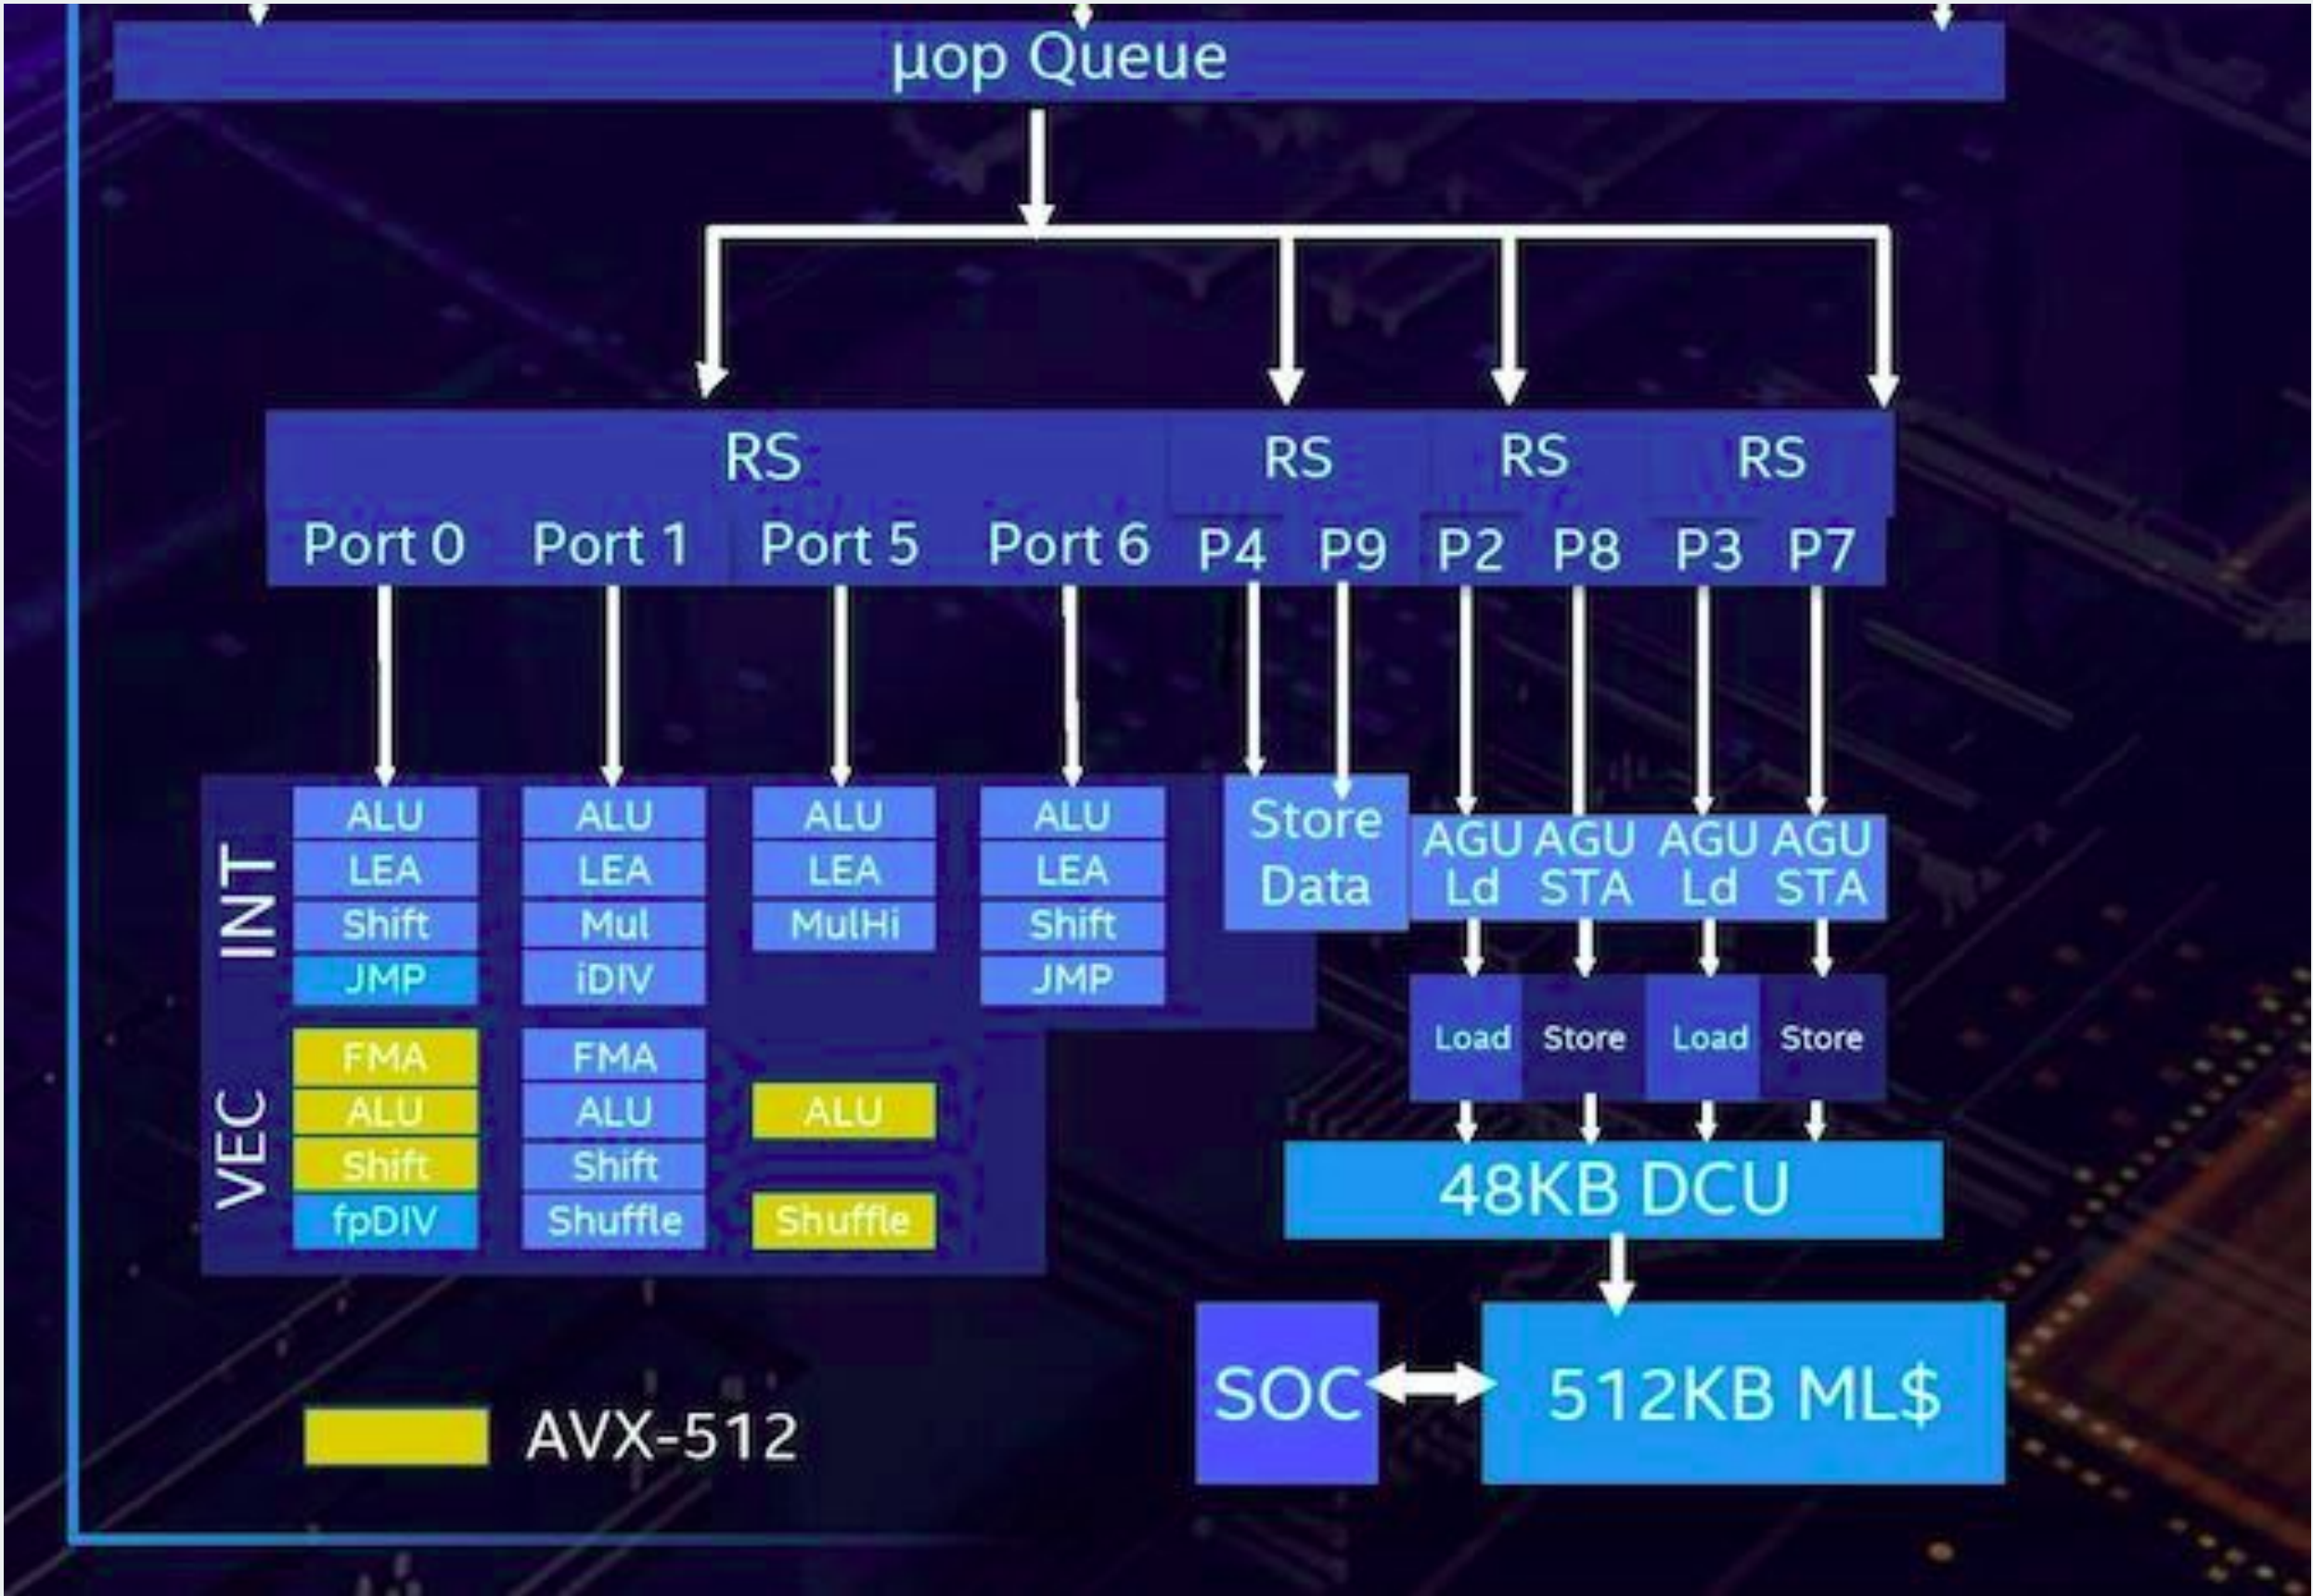


Square requires 528 operations

# RSA Arithmetic Platform Comparison



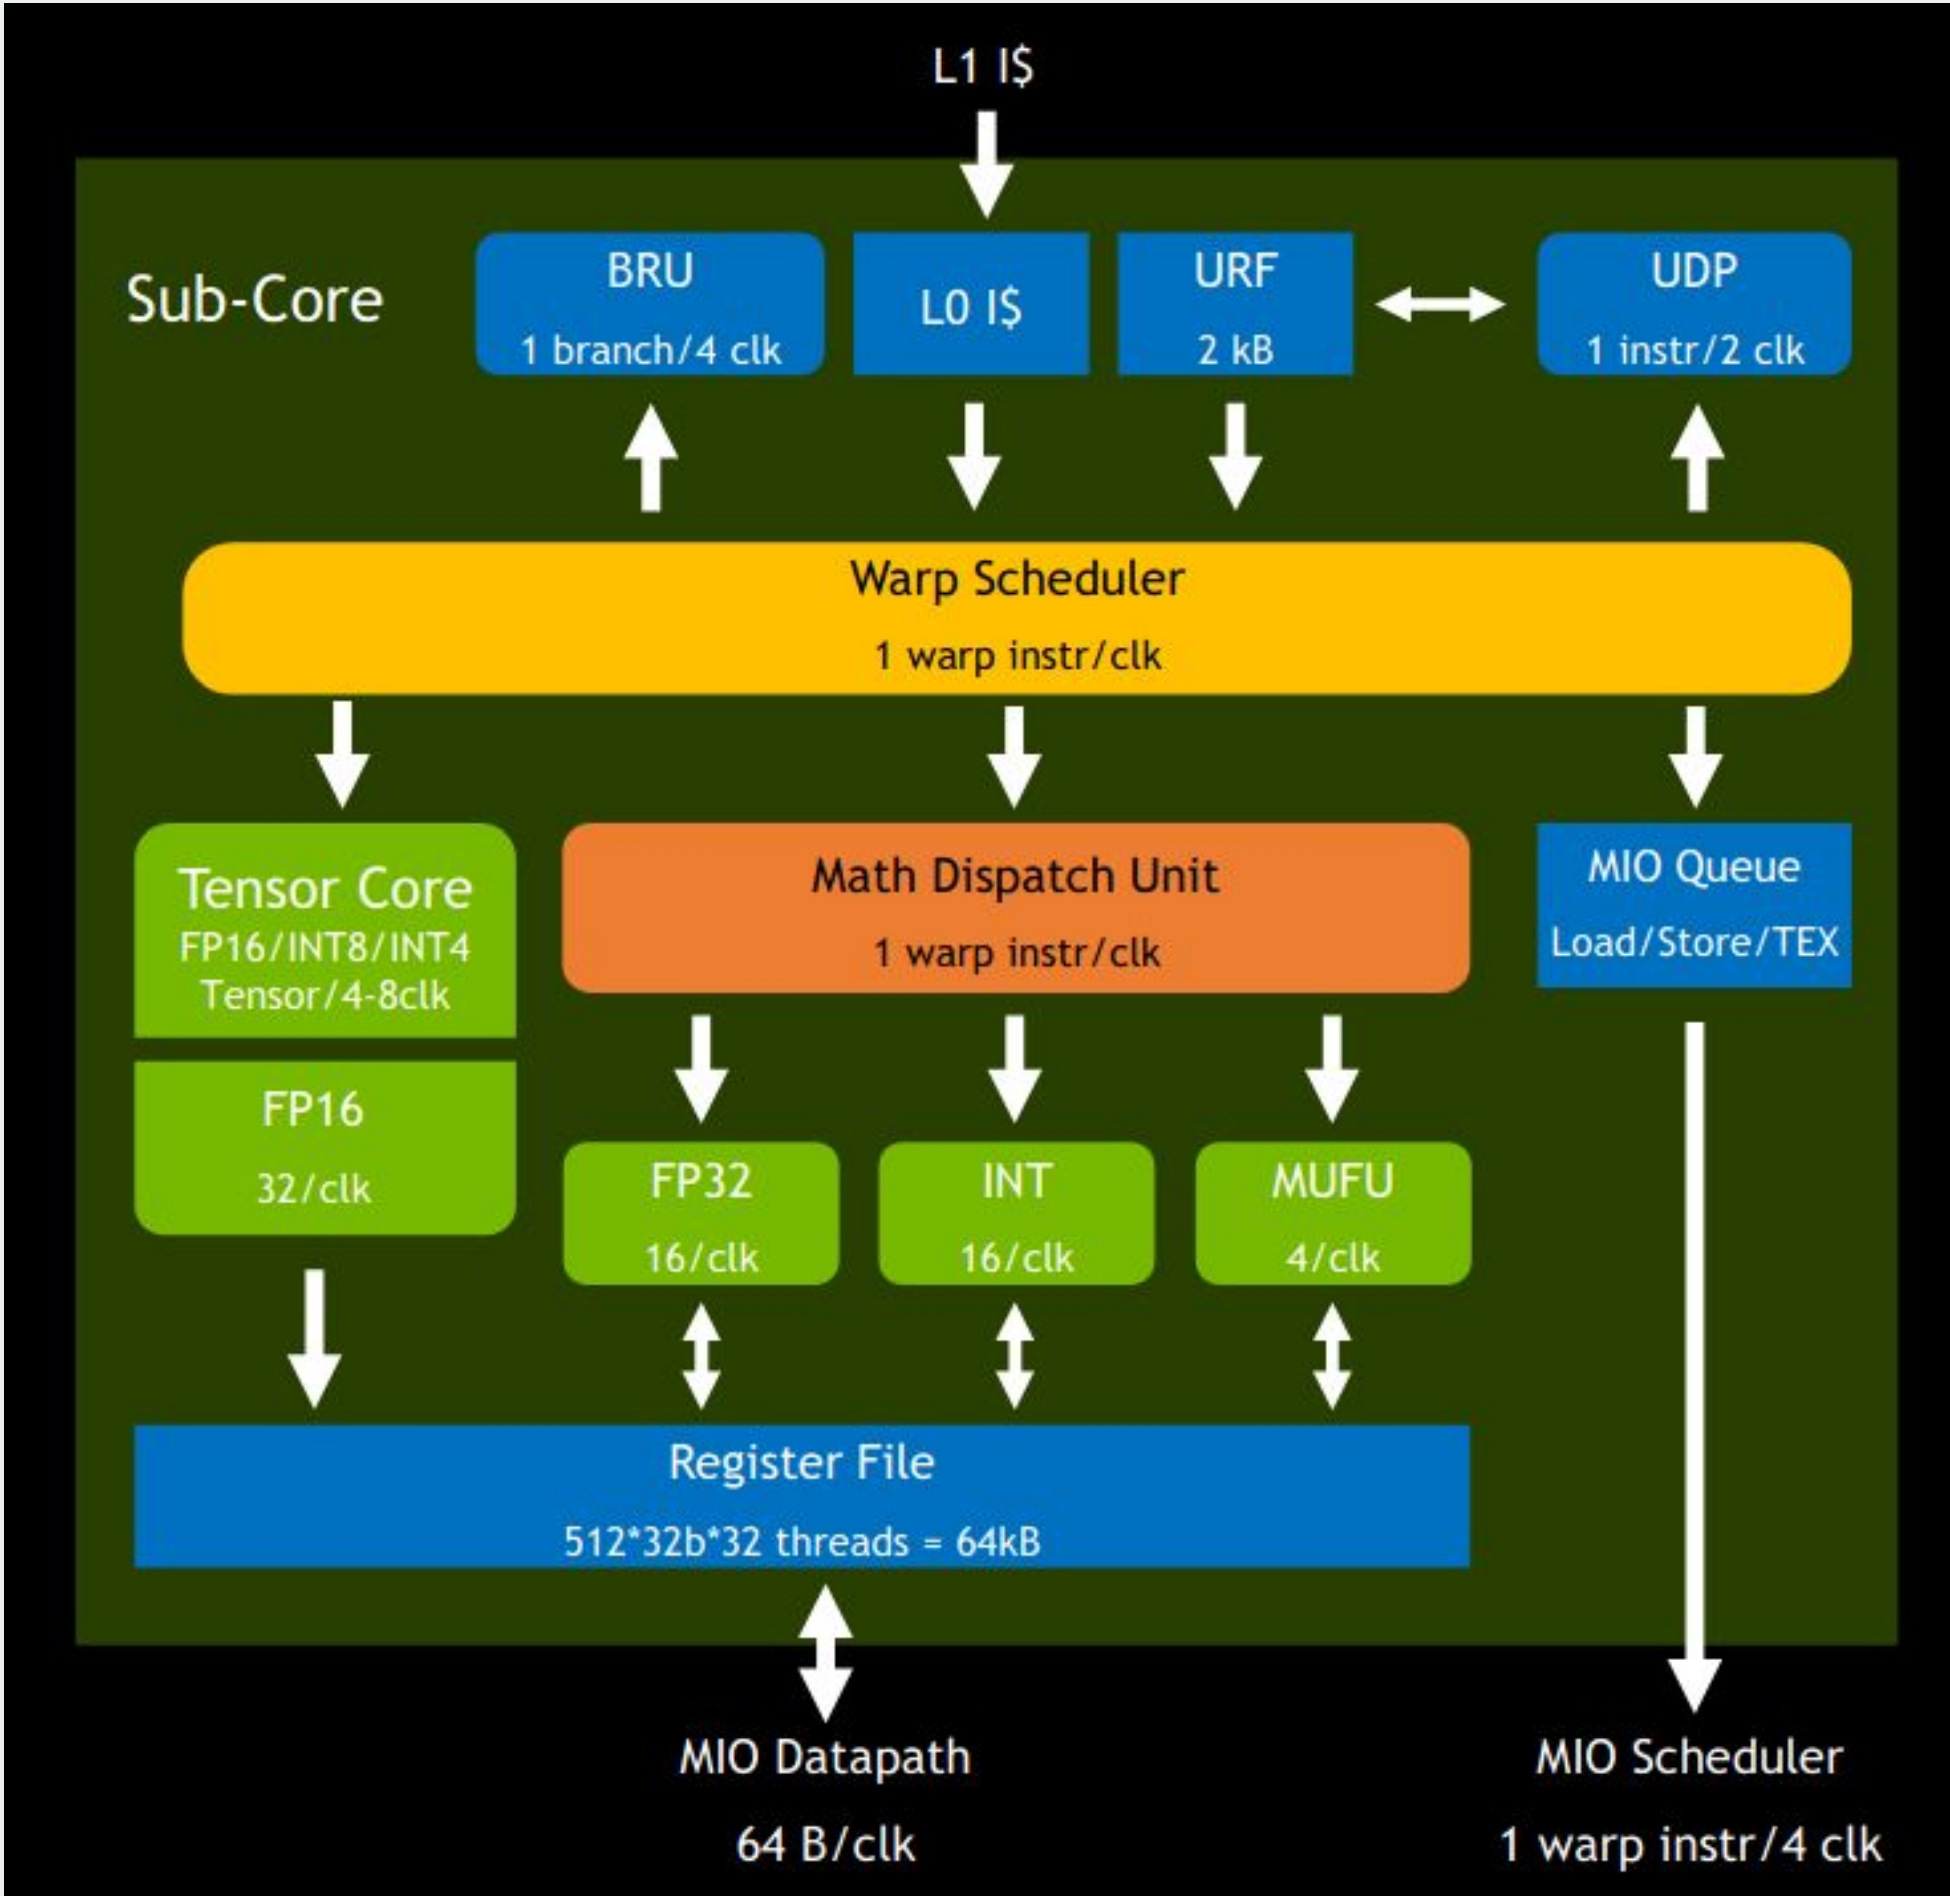
# Performance Impacts - Platform - CPU



Frequency	5 Ghz
Multiplier Size	64-bit
Cores	8
Multiplies Req'd for 2K	1024
Theoretical Throughput	39 M ops/s

2048-bit multiply throughput

# Performance Impacts - Platform - GPU



Frequency	1545 MHz
Multiplier Size	32-bit
Cores	4608
Multiplies Req'd for 2K	4096
Theoretical Throughput	1738 M ops/s

2048-bit multiply throughput



# Performance Impacts - Platform - FPGA

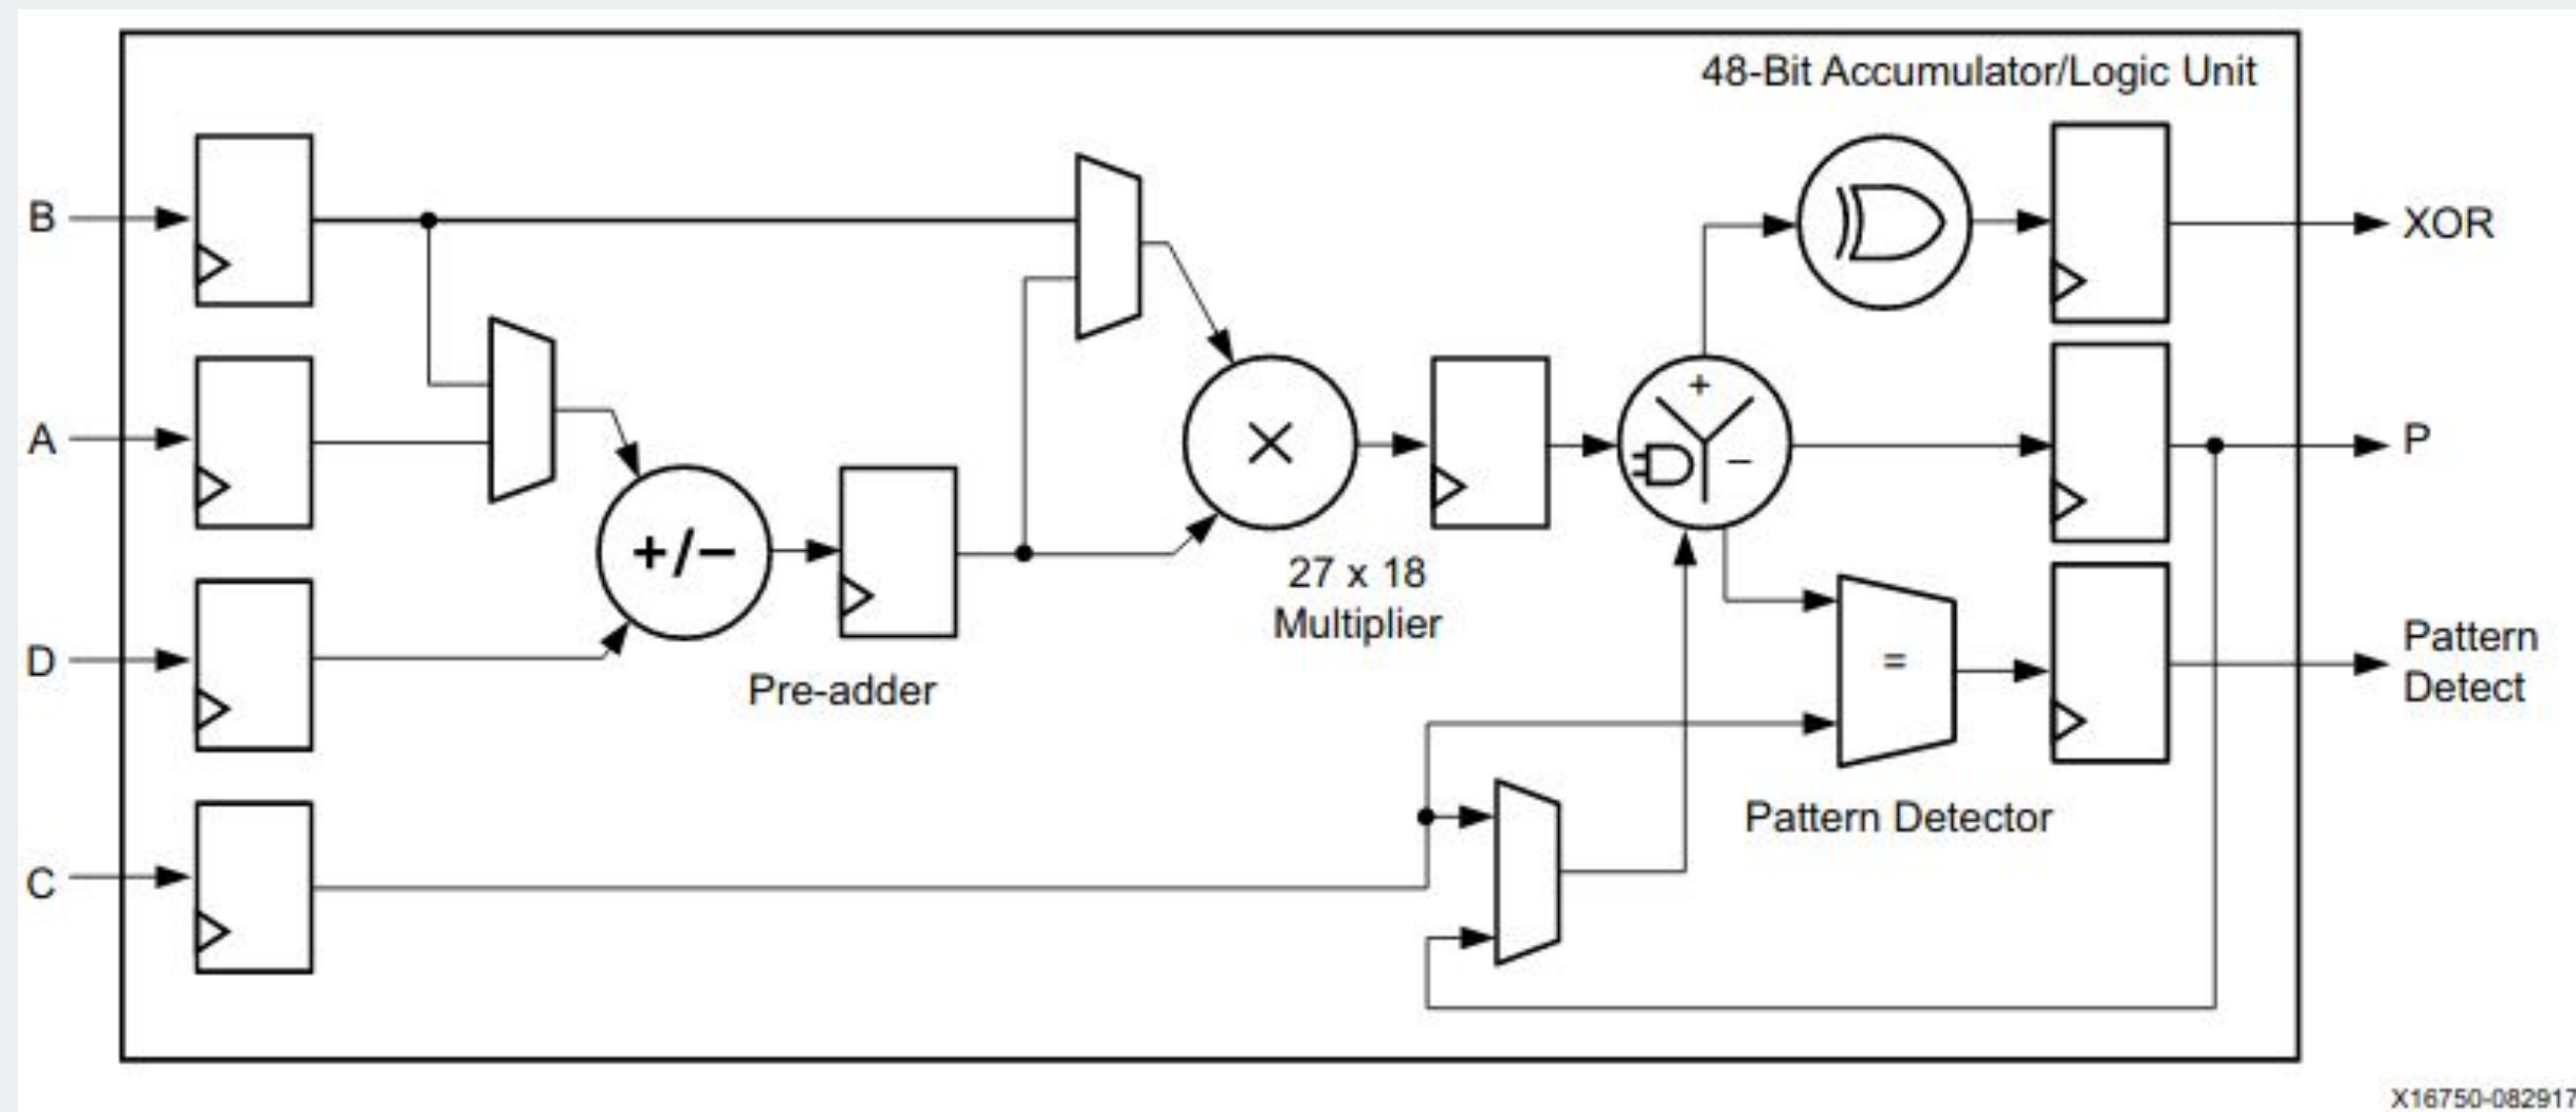



Figure 1-1: Basic DSP48E2 Functionality

Frequency	462 MHz
Multiplier Size	27x18-bit
Cores	6840
Multiplies Req'd for 2K	9678
Theoretical Throughput	327 M ops/s

2048-bit multiply throughput

# Performance Impacts - Platform - ASIC (estimated)



Frequency	1 GHz
Multiplier Size	256-bit
Cores	1,000
Multiplies Req'd for 2K	64
Theoretical Throughput	15,625 M ops/s

2048-bit multiply throughput

Low latency design: 2048 bit redundant representation  
 High throughput design: 256 bit redundant representation

# Performance Impacts - CPU Example

## Scalar

mul-based instruction sequence	mulx-based instruction sequence	mulx/adcx/adox based instruction sequence
<div>mov OP, [pB+8*0]</div> <div>mov rax, [pA+8*0]</div> <div>mul OP</div> <div>add R0, rax</div> <div>adc rdx, 0</div> <div>mov TMP, rdx</div> <div>mov pDst, R0</div> <div>mov rax, [pA+8*1]</div> <div>mul OP</div> <div>mov R0, rdx</div> <div>add R1, rax</div> <div>adc R0, 0</div> <div>add R1, TMP</div> <div>adc R0, 0</div> <div>mov rax, [pA+8*2]</div> <div>mul OP</div> <div>mov TMP, rdx</div> <div>add R2, rax</div> <div>adc TMP, 0</div> <div>add R2, R0</div> <div>adc TMP, 0</div> <div>...</div>	<div>mov OP, [pB+8*0]</div> <div>mulx TMP1,rax, [pA+8*0]</div> <div>add R0, rax</div> <div>adc TMP1, 0</div> <div>mov pDst, R0</div> <div>mulx TMP2,R'0, [pA+8*1]</div> <div>add R'0, R1</div> <div>adc TMP2, 0</div> <div>add R'0, TMP1</div> <div>adc TMP2, 0</div> <div>mulx TMP1,R'1, [pA+8*2]</div> <div>add R'1, R2</div> <div>adc TMP1, 0</div> <div>add R'1, TMP2</div> <div>adc TMP1, 0</div> <div>...</div>	<div>xor rax, rax</div> <div>mov rdx, [pB+8*0]</div> <div>mulx T1, T2, [pA+8*0]</div> <div>adox R0, T2</div> <div>adcx R1, T1</div> <div>mov pDst, R0</div> <div>mulx T1, R'0, [pA+8*1]</div> <div>adox R'0, R1</div> <div>adcx R2, T1</div> <div>mulx T1, R'1, [pA+8*2]</div> <div>adox R'1, R2</div> <div>adcx R3, T1</div> <div>...</div>

## Vector


New instructions:  
512-bit Integer Fused Multiply Add (IFMA)

VPMADD52LUQ zmm1 {k1}{z},  
zmm2, zmm3/m512/m64bcst

VPMADD52HUQ zmm1 {k1}{z},  
zmm2, zmm3/m512/m64bcst


8 lanes of 52x52 + 64  
First instruction does low result accumulate  
Second instruction does high result accumulate

# Throughput






# Throughput




# Latency



# BLS-381 MultiExp Platform Comparison

# BLS-381 MultiExp Throughput



# SNARK ASIC

# SNARK ASIC Strawman Proposal (technical)

- **multiexponentiations only**—no FFTs, no witness generation

# SNARK ASIC Strawman Proposal (technical)

- **multiexponentiations only**—no FFTs, no witness generation
- **BLS12-381 only**—"blockchain standard", secure, efficient recursion (~1m gates overhead)

# SNARK ASIC Strawman Proposal (technical)

- **multiexponentiations only**—no FFTs, no witness generation
- **BLS12-381 only**—"blockchain standard", secure, efficient recursion (~1m gates overhead)
- **28nm process**—no costly FinFET process (sufficient for PCIe4, GDDR5)



# SNARK ASIC Strawman Proposal (technical)

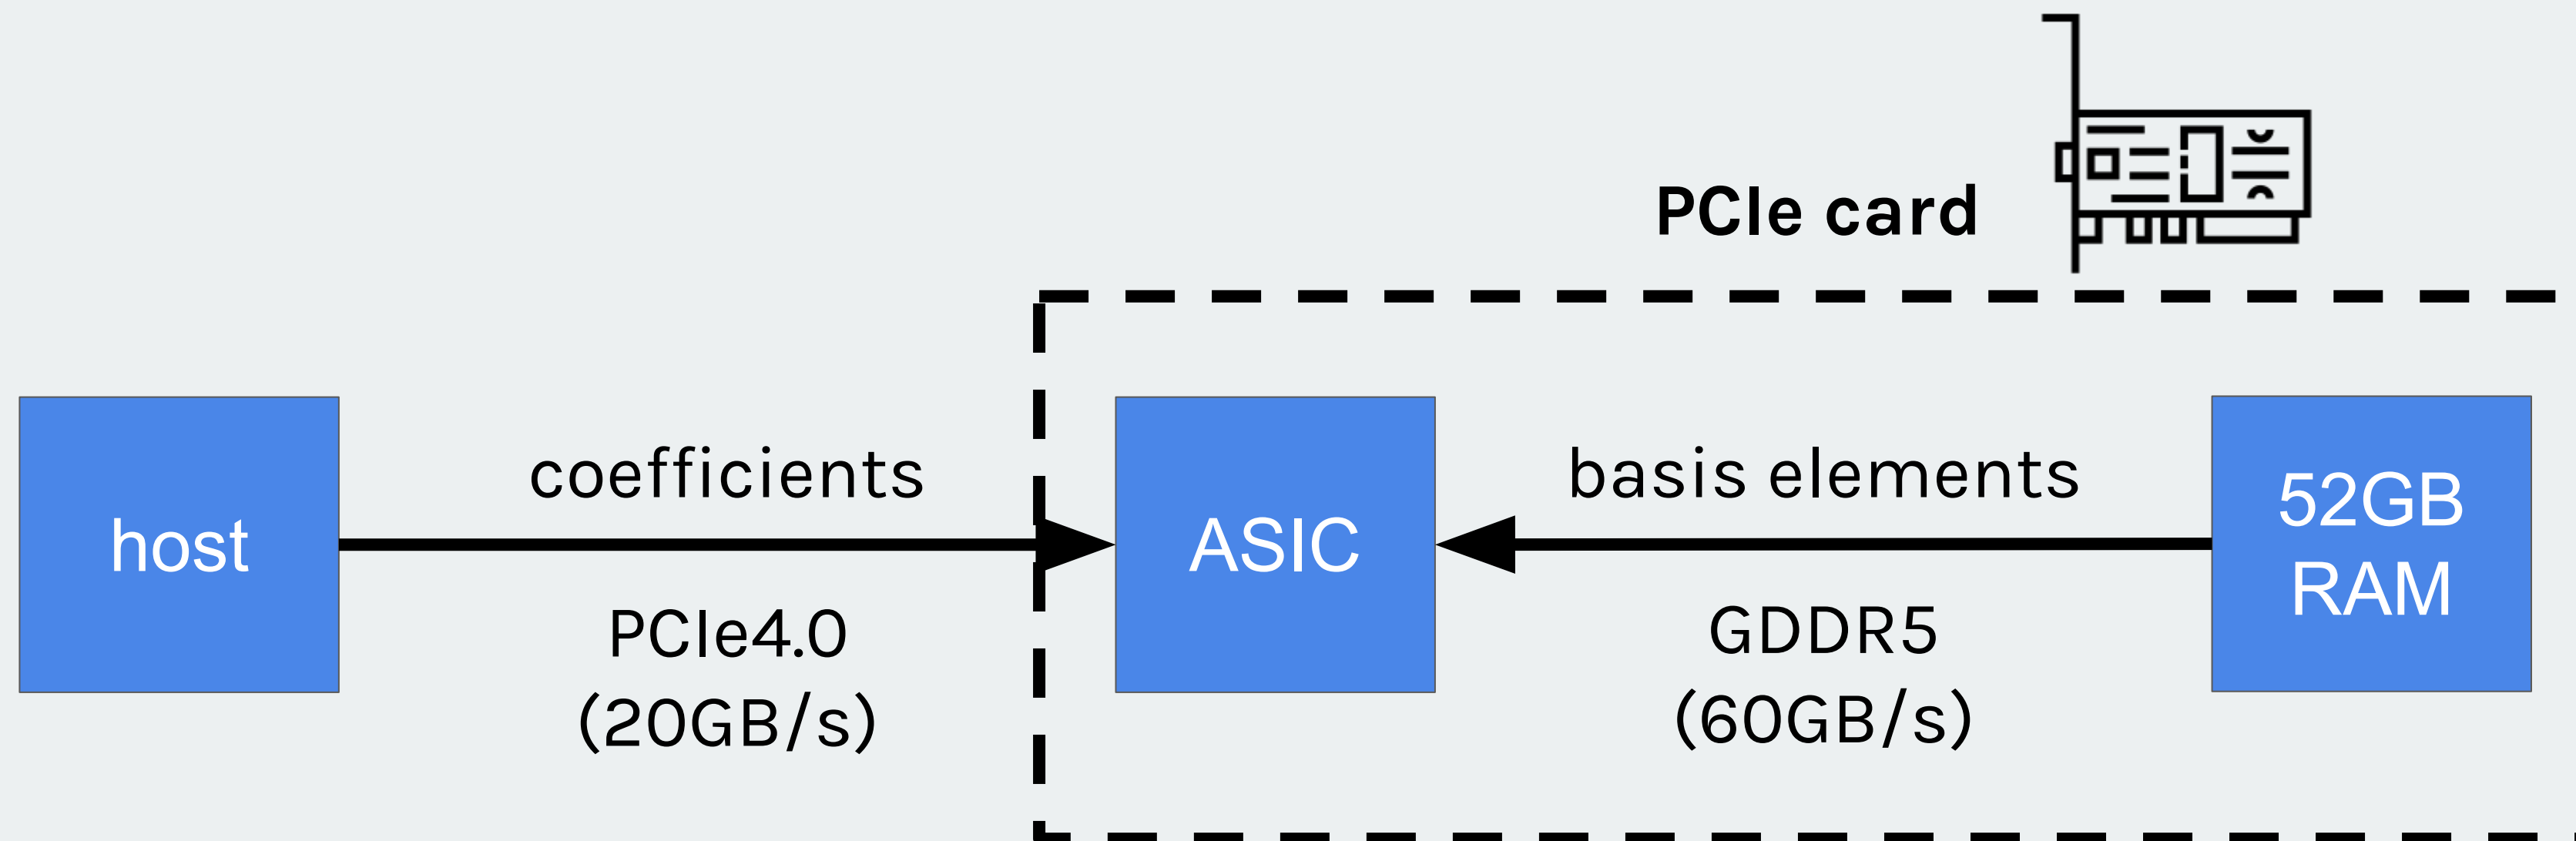
- **multiexponentiations only**—no FFTs, no witness generation
- **BLS12-381 only**—"blockchain standard", secure, efficient recursion (~1m gates overhead)
- **28nm process**—no costly FinFET process (sufficient for PCIe4, GDDR5)
- **PCIe4.0**—cost effective, high bandwidth, convenient

# SNARK ASIC Strawman Proposal (technical)

- **multiexponentiations only**—no FFTs, no witness generation
- **BLS12-381 only**—"blockchain standard", secure, efficient recursion (~1m gates overhead)
- **28nm process**—no costly FinFET process (sufficient for PCIe4, GDDR5)
- **PCIe4.0**—cost effective, high bandwidth, convenient
- **GDDR5 RAM**—4x bandwidth boost versus PCIe alone, less host RAM

# SNARK ASIC Strawman Proposal (technical)

- **multiexponentiations only**—no FFTs, no witness generation
- **BLS12-381 only**—"blockchain standard", secure, efficient recursion (~1m gates overhead)
- **28nm process**—no costly FinFET process (sufficient for PCIe4, GDDR5)
- **PCIe4.0**—cost effective, high bandwidth, convenient
- **GDDR5 RAM**—4x bandwidth boost versus PCIe alone, less host RAM



# **SNARK ASIC Strawman Proposal (non-technical)**

- **Supranational**—trusted to deliver

# **SNARK ASIC Strawman Proposal (non-technical)**

- **Supranational**—trusted to deliver
- **public good**—subsidized NRE (~\$5m), PCIe cards sold at cost (~\$1K)

# **SNARK ASIC Strawman Proposal (non-technical)**

- **Supranational**—trusted to deliver
- **public good**—subsidized NRE (~\$5m), PCIe cards sold at cost (~\$1K)
- **open source**—open source RTL

# **SNARK ASIC Strawman Proposal (non-technical)**

- **Supranational**—trusted to deliver
- **public good**—subsidized NRE (~\$5m), PCIe cards sold at cost (~\$1K)
- **open source**—open source RTL
- **2023**—immediately after the VDF ASIC


# **SNARK ASIC Strawman Proposal (non-technical)**

- **Supranational**—trusted to deliver
- **public good**—subsidized NRE (~\$5m), PCIe cards sold at cost (~\$1K)
- **open source**—open source RTL
- **2023**—immediately after the VDF ASIC
- **funding**—Ethereum Foundation and others



# SNARK ASIC Strawman Proposal (non-technical)

- **Supranational**—trusted to deliver
- **public good**—subsidized NRE (~\$5m), PCIe cards sold at cost (~\$1K)
- **open source**—open source RTL
- **2023**—immediately after the VDF ASIC
- **funding**—Ethereum Foundation and others



VDF Alliance partners

# Open Questions for Discussion

## technical

- **witness generation**—can we optimise to 0.1%? dial-down fancy non-determinism?
- **FFTs**—optional or essential? is 5x incremental speed-up valuable?
- **BLS12-381**—can it become a winner-take-most?
- **bottleneck**—bandwidth or compute? fast multi-point decompression?
- **recursion**—impact to hardware acceleration?

# Open Questions for Discussion

## technical

- **witness generation**—can we optimise to 0.1%? dial-down fancy non-determinism?
- **FFTs**—optional or essential? is 5x incremental speed-up valuable?
- **BLS12-381**—can it become a winner-take-most?
- **bottleneck**—bandwidth or compute? fast multi-point decompression?
- **recursion**—impact to hardware acceleration?

## non-technical

- **adoption**—is 2023 good timing? what are burning use cases?
- **funding**—who's in? (Protocol Labs, Tezos, Interchain, Web3, NSF)



# Questions?

hello@supranational.net

CONFIDENTIAL