



FROMAGER: A scalable toolchain for complex ZK Proofs about software

Alex Malozemoff, *James Parker* (james@galois.com), David Archer
Galois, Inc.

Research focus

Proving properties about software in ZK, efficiently

By properties, we mean:

- Prove **vulnerability** exists without revealing **trigger**
- Prove **safety** of code without revealing its source

By efficiently, we mean:

- Prover **and** verifier efficiency, **minimal** communication
- Handle very large circuits (billions of gates)

**Should you disclose software
vulnerabilities?**

Challenges of vulnerability disclosure



Vendors



Bad actors



Users




Public
Program



Exploit





Fromager


(Cheesecloth + Mac'n'Cheese)



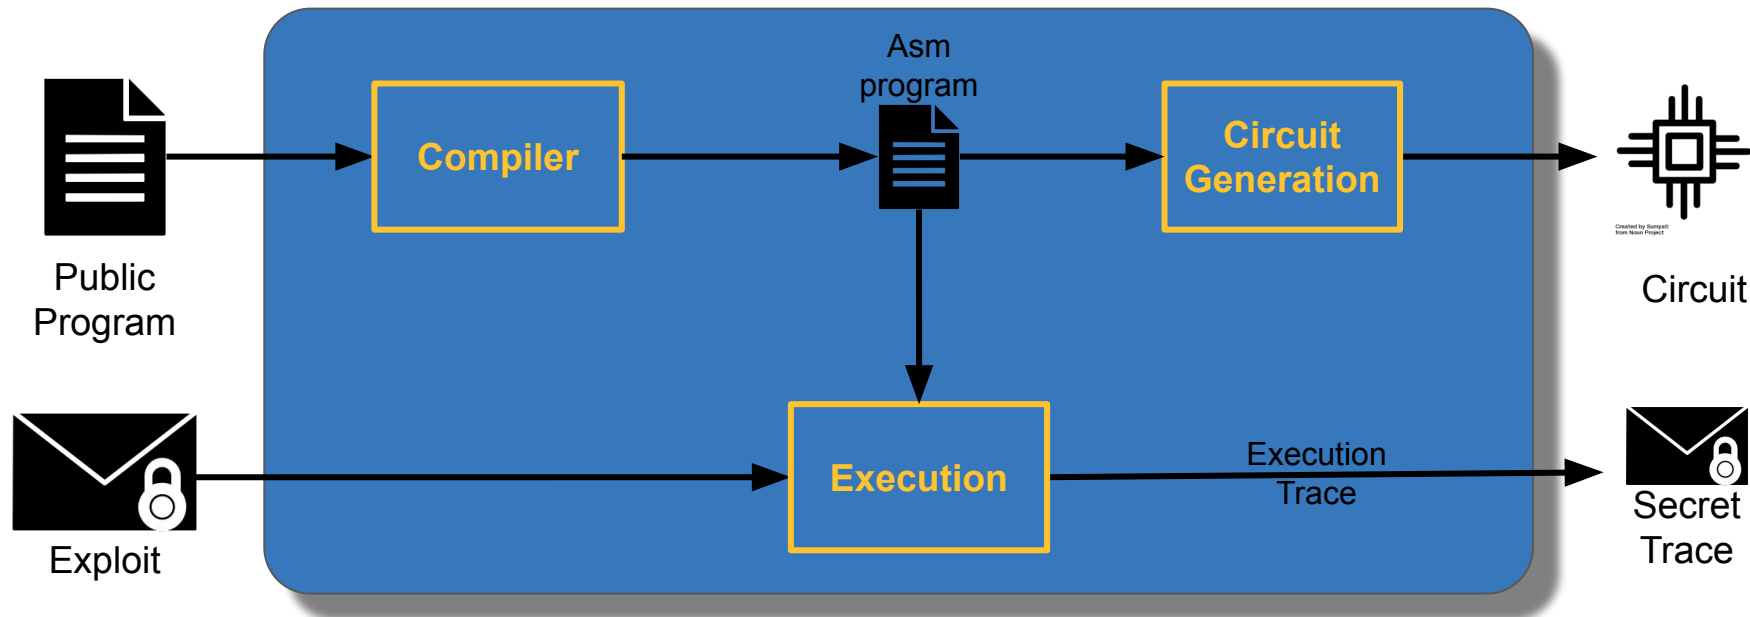
Proof of
vulnerability

Fromager Demonstrations

Vulnerability class	Example Program		Trace length	Gates
Buffer overflow	GRIT		6K	27M
Out-of-bounds array access	FFmpeg (CVE-2013-0864)		80K	673M
Information Leakage	OpenSSL (Heartbleed)		25,000K	17B
Cryptographic protocol bug	Scuttlebutt		WIP	WIP



Cheesecloth - The “Front end”



Created by Samsung
Open Source Project

Circuit



Secret
Trace

Components of the Circuit

- **For each instruction in the program**
 - A circuit that checks whether the CPU state transition was “legal”
- **For each memory reference**
 - A circuit that verifies reads correctly reflect most recent writes
- **For all memory references**
 - A circuit that verifies the ordering used in the above circuit is a correct permutation of the actual execution order

Components of the Circuit

Trace = [state₀, state₁, ..., state_T]
state_i = { registers, program counter, bug_flag }

- State transition function

$$ST(m_i, \text{state}_i) = \text{state}_{i+1}$$

- Memory consistency

$$MC(m_1, m_2, \dots, m_T) = \text{True}$$

State transition function

$$ST(m_i, \text{state}_i) = \text{state}_{i+1}$$

```
fn transition(current_st) -> State {  
  let instr = fetch_instr(current_st.pc);  
  let arg1 = index(current_st.regs, instr.op1);  
  let arg2 = index(current_st.regs, instr.op2);  
  
  let result = mux(instr.opcode == XOR, xor(arg1, arg2),  
                   instr.opcode == ADD, add(arg1, arg2),  
                   ...);  
  ...  
}
```

Memory Consistency Check

$$MC(m_1, m_2, \dots m_T) = \text{True}$$

1. Sort memory operations by address and cycle
 - a. Check that reads match previous write
2. Check that memory operations are a permutation of the execution trace

WRITE	0xdeadbeef	@0	= 42
-------	------------	----	------

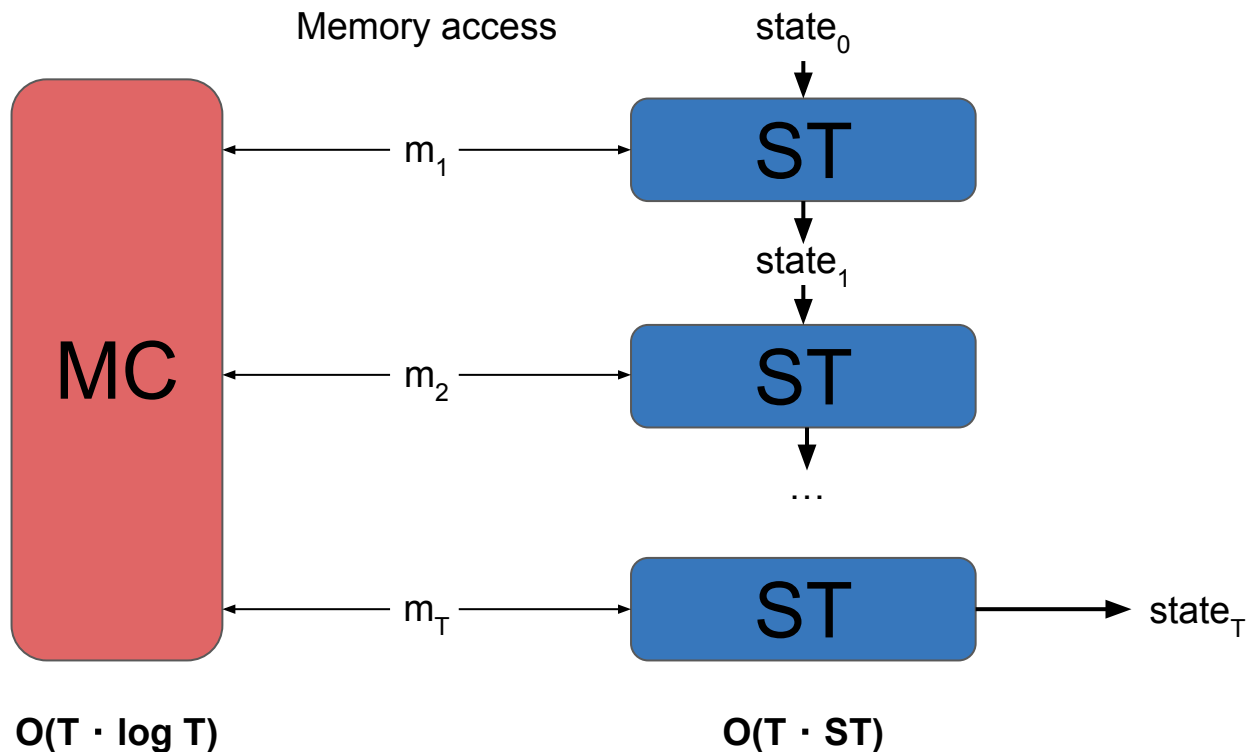
READ	0xdeadbeef	@6	= 42
------	------------	----	------

READ	0xdeadbeef	@9	= 42
------	------------	----	------

WRITE	0xdeadbefe	@2	= 63
-------	------------	----	------

READ	0xdeadbefe	@8	= 63
------	------------	----	------

Circuit generation



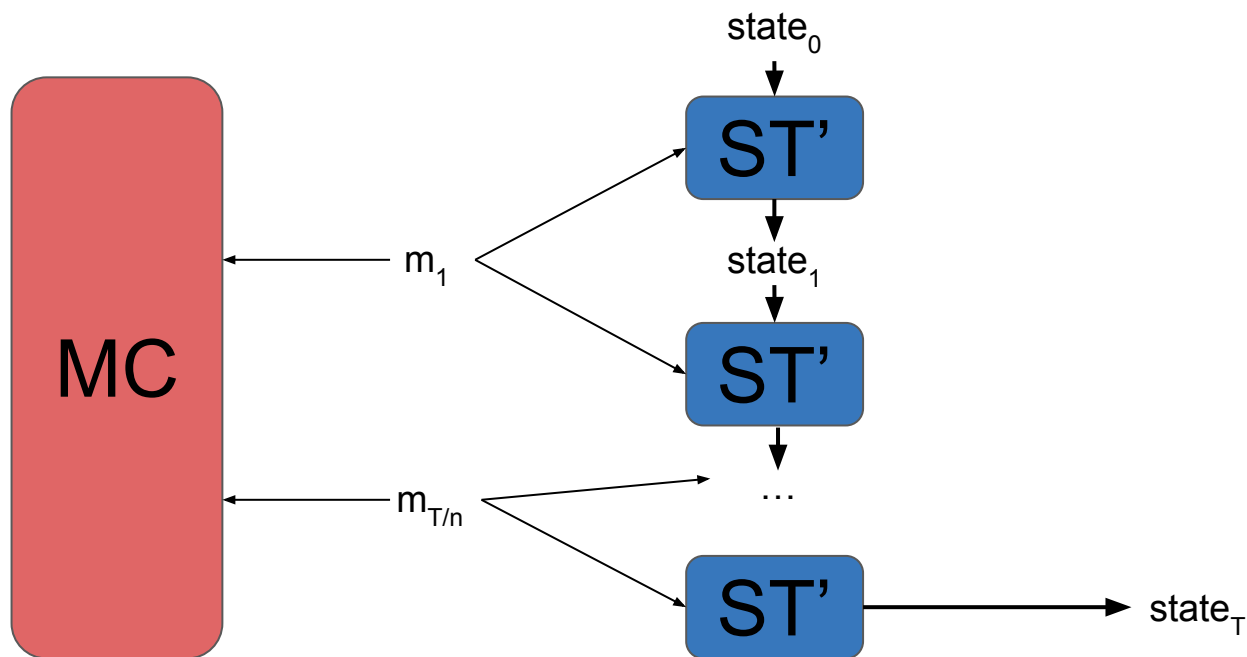
Novel Contributions in Our Framework

- **Optimizations** that minimize circuits that verify program executions
 - a. Sparsity
 - b. Public PC
- Efficient **encodings** to detect vulnerabilities
 - a. Memory errors
 - b. Information leakages

Optimization: Sparsity

- Certain instructions are expensive and used infrequently
 - Example: memory operations
- What if we don't run these instructions every cycle?
- Savings: no need to generate memory check circuits for excluded cycles
 - Instructions that are sparsely used are only inserted once every n cycles.
 - The circuit cost becomes $1 / n$

Optimization: Sparsity



Optimization: Public PC

- Encoding the state transition function for each instruction is expensive
- What if we generate circuits for frequently used basic blocks in the program?
- Savings: The program counter (PC) is public, so we can generate optimized circuit representations for the entire basic block

Optimization: Public PC


```
fn transition(current_st) -> State {  
  let instr = fetch_instr(current_st.pc);  
  let arg1 = index(current_st.regs, instr.op1);  
  let arg2 = index(current_st.regs, instr.op2);  
  
  let result = mux(instr.opcode == XOR, xor(arg1, arg2),  
                   instr.opcode == ADD, add(arg1, arg2),  
                   ...);  
  
  ...  
}
```

Optimization: Public PC

```
fn transition(current_st) -> State {  
  let instr = fetch_instr(current_st.pc);  
  let arg1 = index(current_st.regs, instr.op1);  
  let arg2 = index(current_st.regs, instr.op2);  
  
  let result = mux(instr.opcode == XOR, xor(arg1, arg2),  
               instr.opcode == ADD, add(arg1, arg2),  
               ...);  
  ...  
}
```

Optimization: Public PC

Private execution




Code block

MOV R0 R1

ADD R0 R0 R2

MUL R0 R0 R3

Public execution



Encoding: Memory vulnerabilities

Catch memory vulnerabilities:

- Reads/writes before allocation
- Out-of-bounds access
- Reads/writes after free
- Use after free

Example vulnerabilities proven so far:

- **GRIT**: Gameboy advanced Image Transmogrifier
- **FFmpeg**: Library for handling multimedia



Encoding: Memory vulnerabilities

Memory allocation

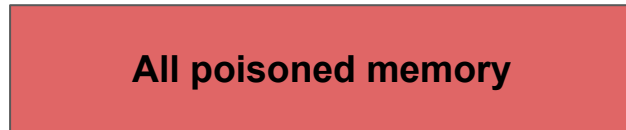
```
ptr = malloc(15);
```



Access to poisoned memory reveal a vulnerability

Memory free

```
free(ptr);
```



Encoding: Memory vulnerabilities

Memory allocation

```
ptr = malloc(15);
```



Memory free

```
free(ptr);
```



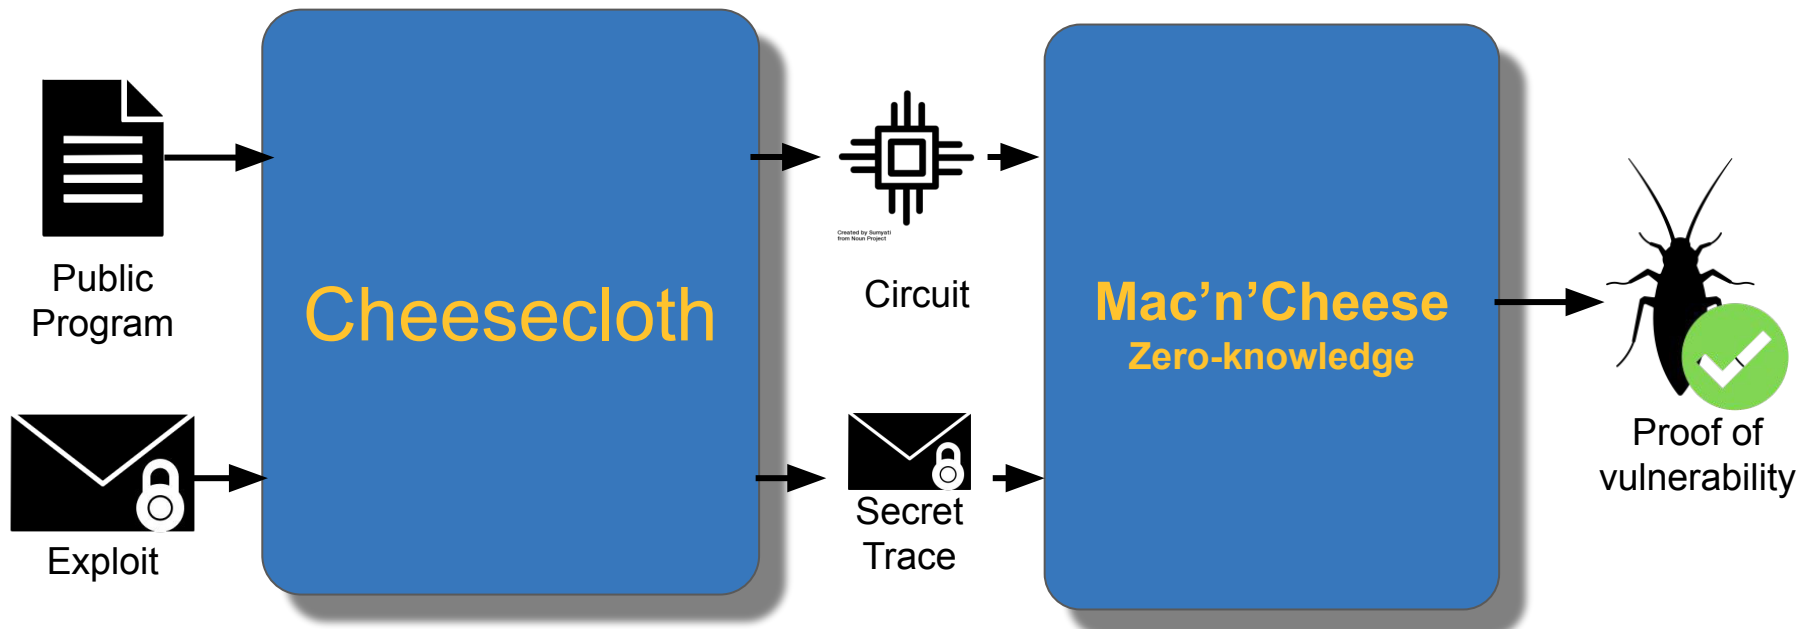
Encoding: Information leakage



Two options:

- **Two executions:** executions should look identical to attacker even when sensitive data changes.
 - Circuit twice as large!
- **Taint tracking:** Mark sensitive data and show that tainted data is revealed to the attacker.
 - Under approximation is non-trivial!

Mac'n'Cheese Backend

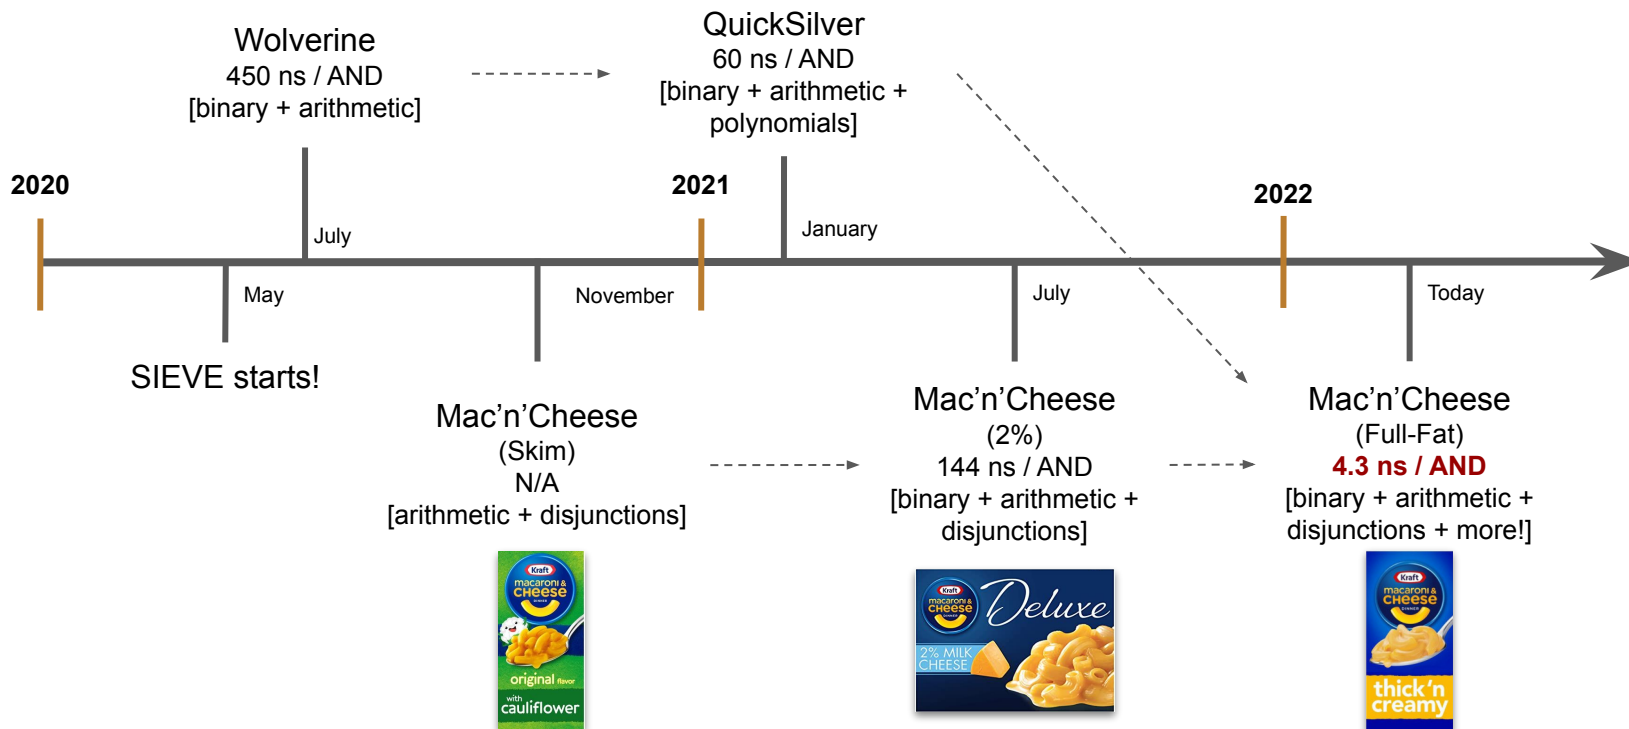


What is Mac'n'Cheese

ZK backend for VOLE-based protocols

- Combines multiple VOLE-ZK protocols in a highly optimized architecture
- Currently: QuickSilver [YSWW21] (for non-disjunctive computations), Mac'n'Cheese (for disjunctive computations) [BMRS21]
- In the pipeline: Appenzeller2Brie [BBMRS21] (for field switching), Mozzarella [BBMS22] (for ring support), AntMan [WYYXW22] (for optimizing for-loops), ...

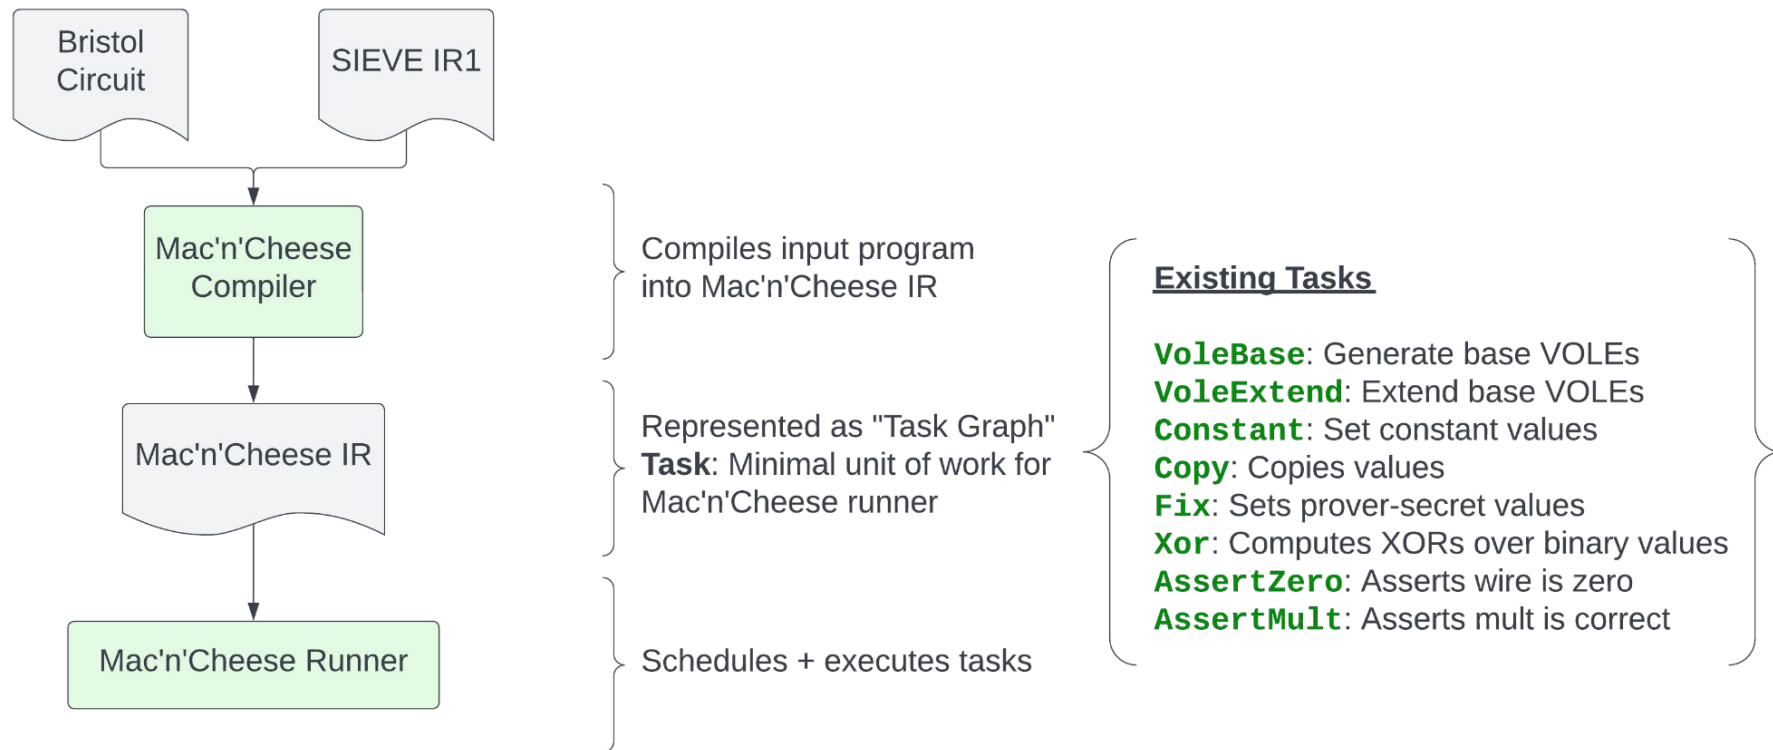
VOLE-ZK implementation timeline



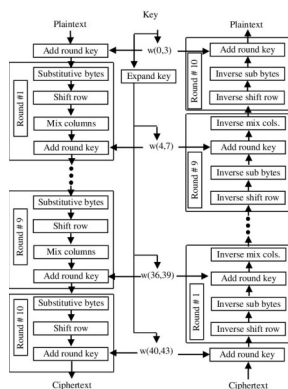
Reference herein to any specific commercial product, process, or service by trade name, trademark or other trade name, manufacturer or otherwise, does not necessarily constitute or imply endorsement by DARPA, the Defense Department or the U.S. government, and shall not be used for advertising or product endorsement purposes.

Sources: https://target.scene7.com/is/image/Target/GUEST_5824921a-638b-4936-b57a-bab60b774594, https://target.scene7.com/is/image/Target/GUEST_0ad8e33a-65b1-49ec-9fb0-a9d7fefed37b, https://target.scene7.com/is/image/Target/GUEST_1dc489ef-a819-4de3-b299-4f258ea28a68

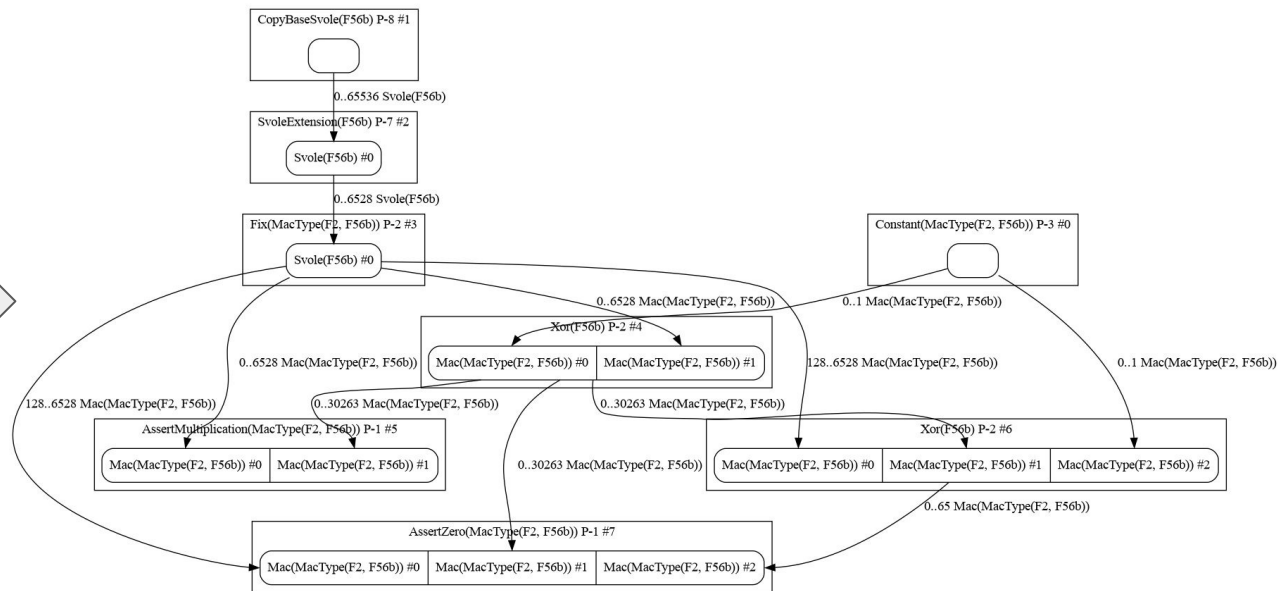
Mac'n'Cheese Architecture



Compiling AES using Mac'n'Cheese



Compilation



Input circuit

Mac'n'Cheese IR

Benefits: Modularity, Parallelizability, Performance

Modular

- Can add tasks easily + independently
 - Ligerio integration, RAM, polynomial checks, etc.
- Can build complex components from tasks
 - Field switching is “just” composition of tasks

Parallelizable

- Versus “sequential” nature of circuit representation

Performant

- Tasks do “one thing”, which can be done much faster

Mac'n'Cheese performance

Experiment:

- Ran on **m5.8xlarge** instances between Virginia and Oregon
- Circuit contained **2.3 B** AND gates + **10 B** XOR gates
- Overall Time: **9.9 seconds**
 - ⇒ **4.3 ns / AND gate**
 - ⇒ **232 million AND gates / second**

Mac'n'Cheese vs software AES + SHA-2

Software AES: **483 ns** per block

- Using <https://github.com/kokke/tiny-AES-c>

Mac'n'Cheese: **51200 ns** per block

⇒ **~106x** slower than software AES!

SHA-2 is similar:

- **~72x** slower versus `hashlib`'s `sha256`!

Mac'n'Cheese run across the US is
~100x slower than native computation*

* For heavily optimized AES / SHA-2 circuits

Availability

- Planning to open source implementation “soon”
 - Reach out if you’re interested in early access!
- Also: if you have use cases that require interactive ZK
 - We’re happy to share access to our implementation / Docker-ized test bed

FROMAGER: A scalable toolchain for complex ZK Proofs about software

| galois |



qedit



Santiago Cuellar



Stuart Pernsteiner



Ben Razet



Bill Harris



Carsten Baum
Peter Scholl



Eran Tromer



Daniel Benarroch



Marc Rosen



Chris Phifer



Dave Archer




Alex Malozemoff

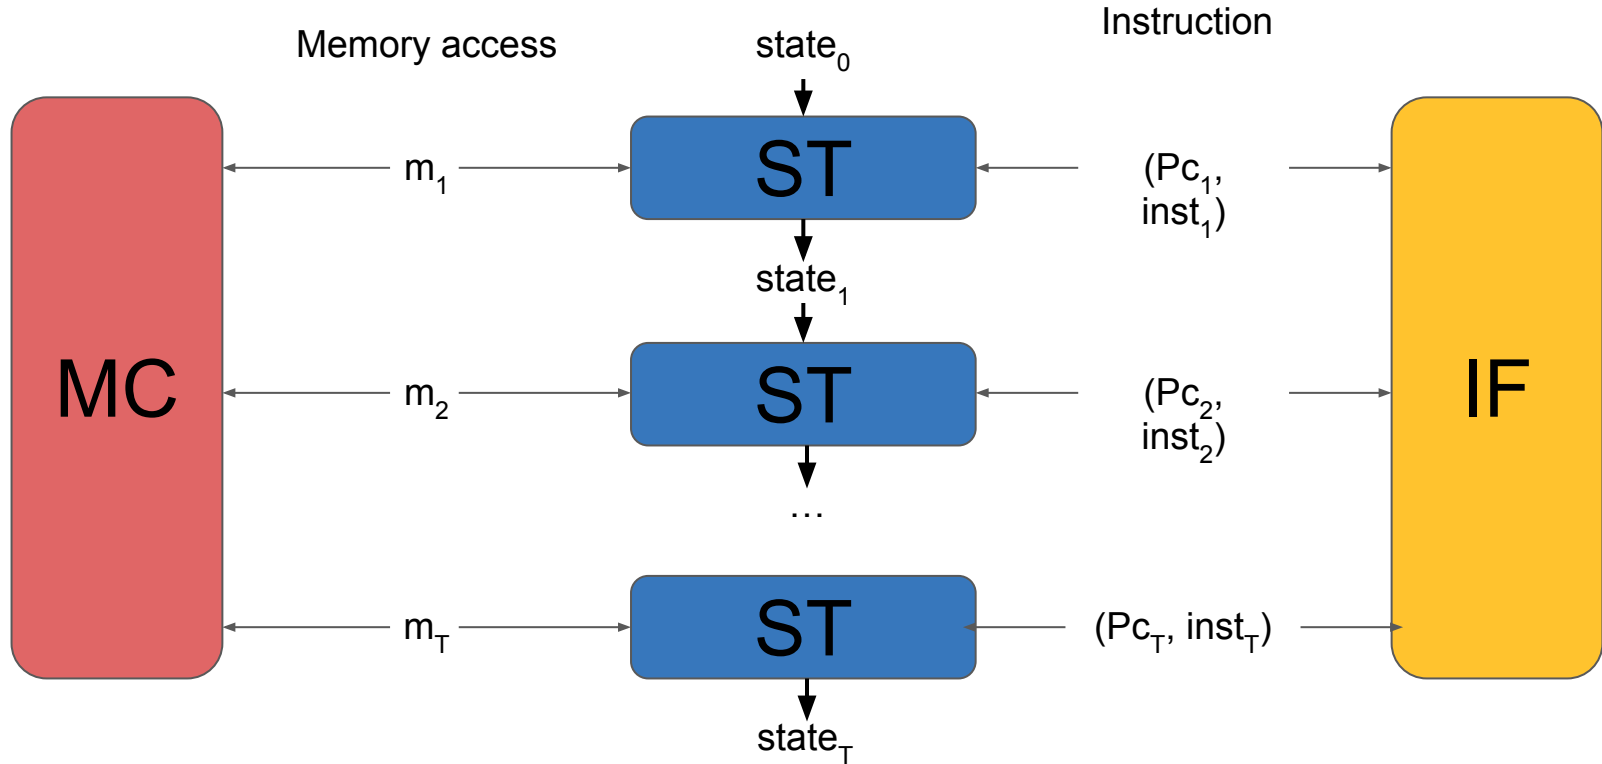


Constance Beguier

Optimization: Public PC



Circuit generation



Execution: Information leakage



Taint tracking

`{secret x}`

`y = x + 1`

Execution: Information leakage



Taint tracking

```
{secret x}  
if x then  
    y = r1;  
else:  
    y = r2;
```

Execution: Information leakage



We use conservative taint tracking (underapproximation):

- Punt on branches and arrays
- Guarantees a vulnerability.

Enough to prove:

- heartbleed
- other protocols that don't branch on secret data.

Encoding: Cryptographic protocol (WIP)



Simulate adversarial program

- Multiple programs communicating through network
- Secret program (adversary)
- Support random-seed
 - Fiat-Shamir doesn't work!

Outline

- 1. Background**
- 2. TA1**
 - a. Intro/Motivation
 - b. Optimizations (sparsity, public pc)
 - c. Vulnerability encodings (memory, taint analysis)
- 3. TA2**
 - a. High level architecture
 - b. Performance numbers
 - c. Open sourcing
- 4. Call for users/use cases**