

The logo consists of the characters 'R' and '0' in a bold, sans-serif font. The 'R' is white and the '0' is black. A diagonal line splits the '0' from the top-left to the bottom-right, with the top-left half being white and the bottom-right half being black.

ZIRGEN: MLIR-based compiler for zk-STARKs

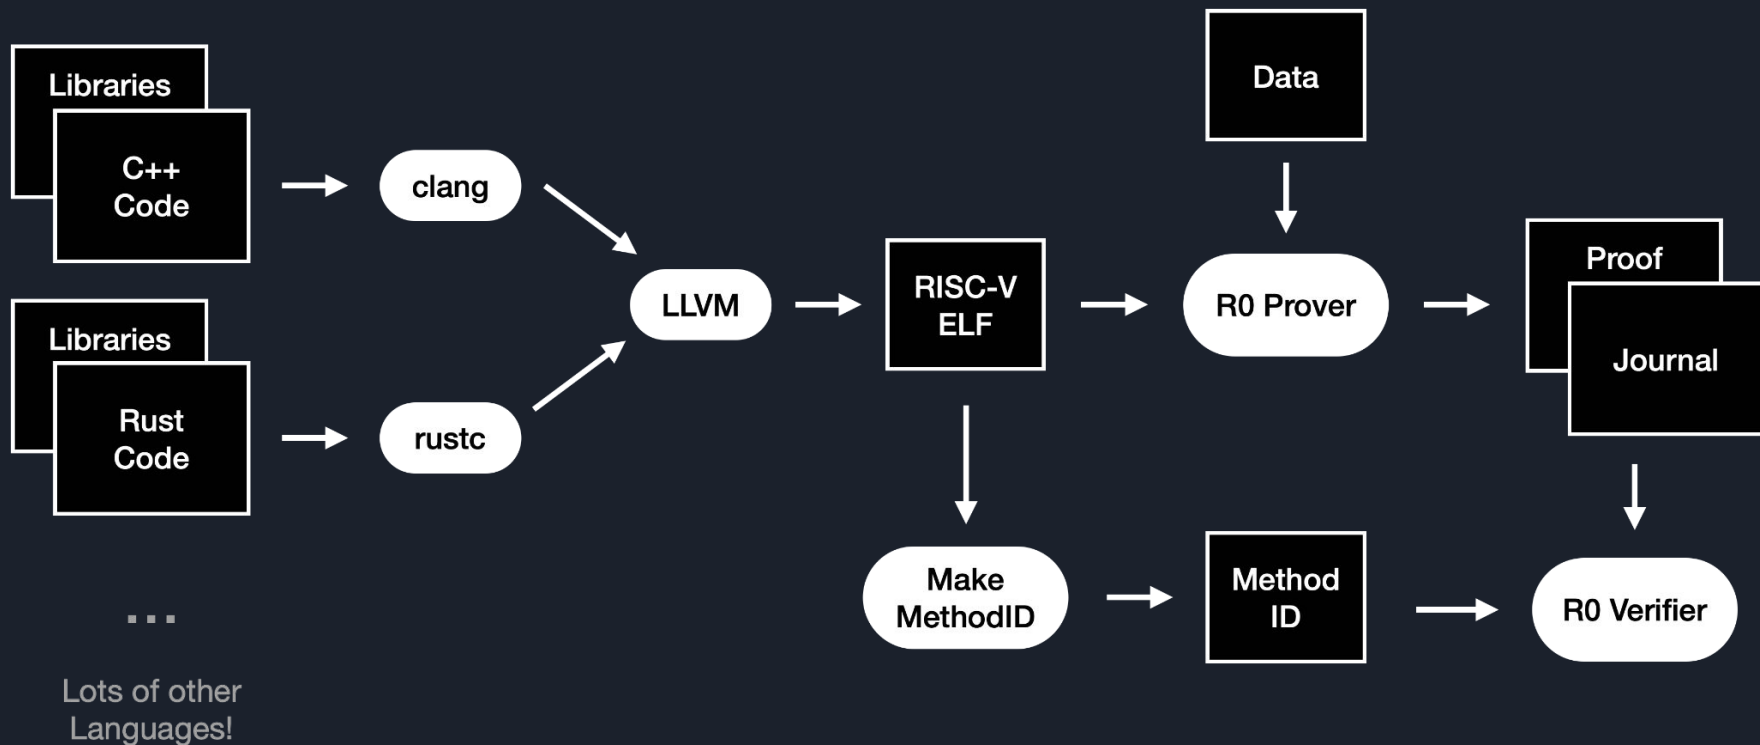
Frank Laub
RISC Zero



What is RISC Zero

- An open source (Apache 2) zkVM, including a working prover and verifier
- Implemented via a zk-STARK proof system
- Which emulates the RISC-V processor ISA

ZKP adoption thru traditional developer toolchains





Why ZIRGEN?

- Multiple complex circuits (RISC-V, SHA256, FFPU)
- Rapidly changing proving technology (Plookup, Hyperplonk, ...)
- Optimizations at different levels
- Multiple hardware backends (CPU, GPU, FPGA, ASICs)
- Multiple abstractions for HW (CUDA, Metal, WebGPU, OpenCL, ...)

^^^

Sounds like a we need a compiler!



What is ZIRGEN?

- C++-based eDSL for constructing arithmetic circuits
- MLIR stack for optimizing circuits and generating constraints
- Hardware-specific code generation
- Support for STARKs, PLONK, and Plookup



ZK Circuits

- Applications
 - Fibonacci
 - Verifier (recursion)*
- Virtual Machines
 - RISC-V
 - MIPS*
 - WASM*
- Accelerators
 - SHA-256
 - Poseidon*
 - BigInt*
 - FFPU (Finite field processing unit)

* Not yet implemented



Compiler approach to circuit generation

- Frontends
 - eDSL
 - DSL
 - Translators
- Passes
 - Canonicalization
 - CSE (Common Subexpression Elimination)
 - DCE (Dead Code Elimination)
 - MakePolynomial
 - ComputeDegree
- Backends
 - C++
 - Rust
 - CUDA
 - FFPU
 - WASM



Compiler pipeline

- fib.cpp (eDSL)
- Zirgen dialect
- Generated source code
 - Step functions
 - step_exec.cpp
 - step_verify_bytes.cpp
 - step_verify_mem.cpp
 - step_accum.cpp
 - Polynomials
 - poly_fp.cpp
 - poly_ext.rs



MLIR: Multi-Level Intermediate Representation

“A collection of modular and reusable software components that enables the progressive lowering of operations, to efficiently target hardware in a common way.”

Google Brain - MLIR team



MLIR: Multi-Level Intermediate Representation

- LLVM IR v2: Lessons learned
- IR matches level of abstraction
- Enables transformations
- Optimizations at appropriate abstraction level
- LIT: LLVM Integrated Tester
- Development workflow
- TableGen

```
def EqualZeroOp : ZirconOp<"eqz", [IsEval]> {  
  let summary = "Require a number to be equal to zero, or fail;"  
  let arguments = (ins Val:$in);  
  let assemblyFormat = [{ $in `:` ` type($in) attr-dict }];  
  let hasCanonicalizer = 1;  
}
```



eDSL: embedded Domain Specific Language

- Module: container for components
- Component: like a function, has inputs/outputs and a body
- Val: generic value type, parameters: field prime, extension size
- Register: a named value
- Buffer: an array of values
- NONDET: Non-deterministic block
- IF: Conditional block



Zirgen Dialect

- Types
 - Val
 - Constraint
 - Buffer
- Attributes
 - Tap
 - Polynomial
- Ops
 - ConstOp
 - NondetOp
 - IfOp
 - AllocOp
 - BackOp
 - SliceOp
 - GetOp
 - SetOp
 - GetGlobalOp
 - SetGlobalOp
- Ops (cont.)
 - EqualZeroOp
 - BarrierOp
 - UnaryOp
 - BinaryOp
 - IsZeroOp
 - BitAndOp
 - FailOp
 - TrueOp
 - AndEqzOp
 - AndCondOp
 - ExternOp



eDSL: C++ example

Fibonacci sequence

```
Module module;
module.addFunc<3>("fib", {const_buf(3), global_buf(1), mut_buf(1)},
  [](Buffer ctrl, Buffer out, Buffer data) {
    Register val = data[0];
    IF(ctrl[0]) { // init
      val = 1;
    }
    IF(ctrl[1]) { // body
      val = BACK(1, val) + BACK(2, val);
    }
    IF(ctrl[2]) { // fini
      out[0] = val;
    }
  });
```



zirgen dialect

```
func.func @fib(%ctrl: !zirgen.buffer<3, constant>,  
              %out: !zirgen.buffer<1, global>,  
              %data: !zirgen.buffer<1, mutable>) ->  
    !zirgen.val<default> {  
    %0 = zirgen.const 1  
    %1 = zirgen.get %ctrl[0] back 0 : <3, constant>  
    zirgen.if %1 : <default> {  
        zirgen.set %data : <1, mutable>[0] = %0 : <default>  
    }  
    %2 = zirgen.get %ctrl[1] back 0 : <3, constant>  
    zirgen.if %2 : <default> {  
        %7 = zirgen.get %data[0] back 1 : <1, mutable>  
        %8 = zirgen.get %data[0] back 2 : <1, mutable>  
        %9 = zirgen.add %7 : <default>, %8 : <default>  
        zirgen.set %data : <1, mutable>[0] = %9 : <default>  
    }  
    %3 = zirgen.get %ctrl[2] back 0 : <3, constant>  
    zirgen.if %3 : <default> {  
        %7 = zirgen.get %data[0] back 0 : <1, mutable>  
        zirgen.set_global %out : <1, global>[0] = %7 : <default>  
    }  
    %4 = zirgen.sub %0 : <default>, %3 : <default>  
    %5 = zirgen.add %1 : <default>, %2 : <default>  
    %6 = zirgen.add %5 : <default>, %3 : <default>  
    return %4 : !zirgen.val<default>  
}
```

Pass: MakePolynomial

Translate execution step function into a polynomial constraint

```
func.func @fib(%arg0: !zirgen.buffer<3, constant>,  
              %arg1: !zirgen.buffer<4, global>,  
              %arg2: !zirgen.buffer<4, mutable>) ->  
    !zirgen.val<default> {  
    %0 = zirgen.const 1  
    %1 = zirgen.get %arg0[0] back 0 : <3, constant>  
    zirgen.if %1 : <default> {  
        zirgen.set %arg2 : <1, mutable>[0] = %0 : <default>  
    }  
    %2 = zirgen.get %arg0[1] back 0 : <3, constant>  
    zirgen.if %2 : <default> {  
        %7 = zirgen.get %arg2[0] back 1 : <1, mutable>  
        %8 = zirgen.get %arg2[0] back 2 : <1, mutable>  
        %9 = zirgen.add %7 : <default>, %8 : <default>  
        zirgen.set %arg2 : <1, mutable>[0] = %9 : <default>  
    }  
    %3 = zirgen.get %arg0[2] back 0 : <3, constant>  
    zirgen.if %3 : <default> {  
        %7 = zirgen.get %arg2[0] back 0 : <1, mutable>  
        zirgen.set_global %arg1 : <1, global>[0] = %7 : <default>  
    }  
    %4 = zirgen.sub %0 : <default>, %3 : <default>  
    %5 = zirgen.add %1 : <default>, %2 : <default>  
    %6 = zirgen.add %5 : <default>, %3 : <default>  
    return %4 : !zirgen.val<default>  
}
```



```
func.func @fib(%arg0: !zirgen.buffer<3, constant>,  
              %arg1: !zirgen.buffer<4, global>,  
              %arg2: !zirgen.buffer<4, mutable>) ->  
    !zirgen.constraint {  
    %0 = zirgen.const 1  
    %1 = zirgen.true  
    %2 = zirgen.get %arg0[0] back 0 : <3, constant>  
    %3 = zirgen.get %arg2[0] back 0 : <1, mutable>  
    %4 = zirgen.sub %3 : <default>, %0 : <default>  
    %5 = zirgen.and_eqz %1, %4 : <default>  
    %6 = zirgen.and_cond %1, %2 : <default>, %5  
    %7 = zirgen.get %arg0[1] back 0 : <3, constant>  
    %8 = zirgen.get %arg2[0] back 1 : <1, mutable>  
    %9 = zirgen.get %arg2[0] back 2 : <1, mutable>  
    %10 = zirgen.add %8 : <default>, %9 : <default>  
    %11 = zirgen.sub %3 : <default>, %10 : <default>  
    %12 = zirgen.and_eqz %1, %11 : <default>  
    %13 = zirgen.and_cond %6, %7 : <default>, %12  
    %14 = zirgen.get %arg0[2] back 0 : <3, constant>  
    %15 = zirgen.get_global %arg1[0] : <1, global>  
    %16 = zirgen.sub %15 : <default>, %3 : <default>  
    %17 = zirgen.and_eqz %1, %16 : <default>  
    %18 = zirgen.and_cond %13, %14 : <default>, %17  
    return %18 : !zirgen.constraint  
}
```

Pass: Canonicalization/CSE

```
func.func @fib(%arg0: !zirgen.buffer<3, constant>,  
              %arg1: !zirgen.buffer<1, global>,  
              %arg2: !zirgen.buffer<1, mutable>)
```

->

```
    !zirgen.constraint {  
    %0 = zirgen.true  
    %1 = zirgen.const 1  
    %2 = zirgen.get %arg0[0] back 0 : <3, constant>  
    %3 = zirgen.true  
    %4 = zirgen.get %arg2[0] back 0 : <1, mutable>  
    %5 = zirgen.sub %4 : <default>, %1 : <default>  
    %6 = zirgen.and eqz %3, %5 : <default>  
    %7 = zirgen.and cond %0, %2 : <default>, %6  
    %8 = zirgen.get %arg0[1] back 0 : <3, constant>  
    %9 = zirgen.true  
    %10 = zirgen.get %arg2[0] back 1 : <1, mutable>  
    %11 = zirgen.get %arg2[0] back 2 : <1, mutable>  
    %12 = zirgen.add %10 : <default>, %11 : <default>  
    %13 = zirgen.get %arg2[0] back 0 : <1, mutable>  
    %14 = zirgen.sub %13 : <default>, %12 : <default>  
    %15 = zirgen.and eqz %9, %14 : <default>  
    %16 = zirgen.and cond %7, %8 : <default>, %15  
    %17 = zirgen.get %arg0[2] back 0 : <3, constant>  
    %18 = zirgen.true  
    %19 = zirgen.get %arg2[0] back 0 : <1, mutable>  
    %20 = zirgen.get global %arg1[0] : <1, global>  
    %21 = zirgen.sub %20 : <default>, %19 : <default>  
    %22 = zirgen.and eqz %18, %21 : <default>  
    %23 = zirgen.and cond %16, %17 : <default>, %22  
    %24 = zirgen.sub %1 : <default>, %17 : <default>  
    %25 = zirgen.add %2 : <default>, %8 : <default>  
    %26 = zirgen.add %25 : <default>, %17 : <default>  
    return %23 : !zirgen.constraint  
}
```



```
func.func @fib(%arg0: !zirgen.buffer<3, constant>,  
              %arg1: !zirgen.buffer<1, global>,  
              %arg2: !zirgen.buffer<1, mutable>) ->  
    !zirgen.constraint {  
    %0 = zirgen.const 1  
    %1 = zirgen.true  
    %2 = zirgen.get %arg0[0] back 0 : <3, constant>  
    %3 = zirgen.get %arg2[0] back 0 : <1, mutable>  
    %4 = zirgen.sub %3 : <default>, %0 : <default>  
    %5 = zirgen.and eqz %1, %4 : <default>  
    %6 = zirgen.and cond %1, %2 : <default>, %5  
    %7 = zirgen.get %arg0[1] back 0 : <3, constant>  
    %8 = zirgen.get %arg2[0] back 1 : <1, mutable>  
    %9 = zirgen.get %arg2[0] back 2 : <1, mutable>  
    %10 = zirgen.add %8 : <default>, %9 : <default>  
    %11 = zirgen.sub %3 : <default>, %10 : <default>  
    %12 = zirgen.and eqz %1, %11 : <default>  
    %13 = zirgen.and cond %6, %7 : <default>, %12  
    %14 = zirgen.get %arg0[2] back 0 : <3, constant>  
    %15 = zirgen.get global %arg1[0] : <1, global>  
    %16 = zirgen.sub %15 : <default>, %3 : <default>  
    %17 = zirgen.and eqz %1, %16 : <default>  
    %18 = zirgen.and cond %13, %14 : <default>, %17  
    return %18 : !zirgen.constraint  
}
```


Pass: ComputeDegree

Annotate ops with the degree for a given polynomial

```
func.func @fib(%arg0: !zirgen.buffer< 3, constant>,  
              %arg1: !zirgen.buffer< 1, global>,  
              %arg2: !zirgen.buffer< 1, mutable>)  
->  
    !zirgen.constraint {  
    %0 = zirgen.const 1  
    %1 = zirgen.true  
    %2 = zirgen.get %arg0[0] back 0 : <3, constant>  
    %3 = zirgen.get %arg2[0] back 0 : <1, mutable>  
    %4 = zirgen.sub %3 : <default>, %0 : <default>  
    %5 = zirgen.and_eqz %1, %4 : <default>  
    %6 = zirgen.and_cond %1, %2 : <default>, %5  
    %7 = zirgen.get %arg0[1] back 0 : <3, constant>  
    %8 = zirgen.get %arg2[0] back 1 : <1, mutable>  
    %9 = zirgen.get %arg2[0] back 2 : <1, mutable>  
    %10 = zirgen.add %8 : <default>, %9 : <default>  
    %11 = zirgen.sub %3 : <default>, %10 : <default>  
    %12 = zirgen.and_eqz %1, %11 : <default>  
    %13 = zirgen.and_cond %6, %7 : <default>, %12  
    %14 = zirgen.get %arg0[2] back 0 : <3, constant>  
    %15 = zirgen.get_global %arg1[0] : <1, global>  
    %16 = zirgen.sub %15 : <default>, %3 : <default>  
    %17 = zirgen.and_eqz %1, %16 : <default>  
    %18 = zirgen.and_cond %13, %14 : <default>, %17  
    return %18 : !zirgen.constraint  
}
```



```
func.func @fib(%arg0: !zirgen.buffer<3, constant>,  
              %arg1: !zirgen.buffer<1, global>,  
              %arg2: !zirgen.buffer<1, mutable>) ->  
    !zirgen.constraint attributes {deg = 2} {  
    %0 = zirgen.const 1 {deg = 0}  
    %1 = zirgen.true {deg = 0}  
    %2 = zirgen.get %arg0[0] back 0 : <3, constant> {deg = 1}  
    %3 = zirgen.get %arg2[0] back 0 : <1, mutable> {deg = 1}  
    %4 = zirgen.sub %3 : <default>, %0 : <default> {deg = 1}  
    %5 = zirgen.and_eqz %1, %4 : <default> {deg = 1}  
    %6 = zirgen.and_cond %1, %2 : <default>, %5 {deg = 2}  
    %7 = zirgen.get %arg0[1] back 0 : <3, constant> {deg = 1}  
    %8 = zirgen.get %arg2[0] back 1 : <1, mutable> {deg = 1}  
    %9 = zirgen.get %arg2[0] back 2 : <1, mutable> {deg = 1}  
    %10 = zirgen.add %8 : <default>, %9 : <default> {deg = 1}  
    %11 = zirgen.sub %3 : <default>, %10 : <default> {deg = 1}  
    %12 = zirgen.and_eqz %1, %11 : <default> {deg = 1}  
    %13 = zirgen.and_cond %6, %7 : <default>, %12 {deg = 2}  
    %14 = zirgen.get %arg0[2] back 0 : <3, constant> {deg = 1}  
    %15 = zirgen.get_global %arg1[0] : <1, global> {deg = 0}  
    %16 = zirgen.sub %15 : <default>, %3 : <default> {deg = 1}  
    %17 = zirgen.and_eqz %1, %16 : <default> {deg = 1}  
    %18 = zirgen.and_cond %13, %14 : <default>, %17 {deg = 2}  
    return {deg = 2} %18 : !zirgen.constraint  
}
```



Codegen: C++

```
Fp step_exec(size_t steps, size_t cycle, Fp** args) {  
    size_t mask = steps - 1;  
    Fp x0(1);  
    auto x1 = args[0][0 * steps + ((cycle - 0) & mask)];  
    if (x1 != 0) {  
        auto& reg = args[2][0 * steps + cycle];  
        reg = x0;  
    }  
    auto x2 = args[0][1 * steps + ((cycle - 0) & mask)];  
    if (x2 != 0) {  
        auto x3 = args[2][0 * steps + ((cycle - 1) & mask)];  
        auto x4 = args[2][0 * steps + ((cycle - 2) & mask)];  
        auto x5 = x3 + x4;  
        auto& reg = args[2][0 * steps + cycle];  
        reg = x5;  
    }  
    auto x6 = args[0][2 * steps + ((cycle - 0) & mask)];  
    if (x6 != 0) {  
        auto x7 = args[2][0 * steps + ((cycle - 0) & mask)];  
        args[1][0] = x7;  
    }  
    auto x8 = x0 - x6;  
    return x8;  
}
```



Performance

- CPU: 25 Khz (cycles/sec)
- GPU: 50 Khz (cycles/sec)
- EOY target: 500 Khz (cycles/sec)
- Recursion: 37s on NVIDIA RTX A5000



Questions?

frank@risczero.com

We are:

- Hiring
 - <https://risczero.com/careers>
- Hacking
 - <https://github.com/risc0>
- Hanging out
 - @risczero on Twitter
 - <https://discord.gg/risczero>

MLIR: Operation

