

From NAND to Verifiable TETRIS



Omer Shlomovits , ZKProof 5 Workshop

Scope



Scope

Algorithms	Software	Hardware



Scope

Algorithms	Software	Hardware
Groth16		



Scope

Algorithms	Software	Hardware
Groth16	C++ implementation <ul style="list-style-type: none">• We can run G_1, G_2 in parallel	



Scope

Algorithms	Software	Hardware
Groth16	<p>C++ implementation</p> <ul style="list-style-type: none">• We can run G_1, G_2 in parallel	<p>Clock runs at frequency f.</p> <ul style="list-style-type: none">• Here is a circuit that runs G_1 in minimal clock cycles• Here is a machine that runs G_2 under area $A[\text{mm}^2]$• Here is a machine that runs Groth16 taking $<5[\text{Watts}]$



Scope

Algorithms	Software	Hardware
Groth16	<p>C++ implementation</p> <ul style="list-style-type: none">• We can run G_1, G_2 in parallel	<p>Clock runs at frequency f.</p> <ul style="list-style-type: none">• Here is a circuit that runs G_1 in minimal clock cycles• Here is a machine that runs G_2 under area $A[\text{mm}^2]$• Here is a machine that runs Groth16 taking $<5[\text{Watts}]$

- For HW to make sense, we must focus on real-world systems



Scope

Algorithms	Software	Hardware
Groth16	<p>C++ implementation</p> <ul style="list-style-type: none">• We can run G_1, G_2 in parallel	<p>Clock runs at frequency f.</p> <ul style="list-style-type: none">• Here is a circuit that runs G_1 in minimal clock cycles• Here is a machine that runs G_2 under area $A[\text{mm}^2]$• Here is a machine that runs Groth16 taking $<5[\text{Watts}]$

- For HW to make sense, we must focus on real-world systems
- Bias towards SNARKs, STARKs



Scope

Algorithms	Software	Hardware
Groth16	<p>C++ implementation</p> <ul style="list-style-type: none">• We can run G_1, G_2 in parallel	<p>Clock runs at frequency f.</p> <ul style="list-style-type: none">• Here is a circuit that runs G_1 in minimal clock cycles• Here is a machine that runs G_2 under area $A[\text{mm}^2]$• Here is a machine that runs Groth16 taking $<5[\text{Watts}]$

- For HW to make sense, we must focus on real-world systems
- Bias towards SNARKs, STARKs
- Proving time is the main bottleneck



ZK Hardware Landscape

Much action, much stealth



PipeZK: Accelerating Zero-Knowledge Proof with a Pipelined Architecture

Ye Zhang^{1,5}

Shuo Wang¹
Cong Wang⁸

Xian Zhang³
Dong Zhou²

Jiangbin Dong^{4,7}
Mingyu Gao^{2,7*}

Xingzhong Mao⁷
Guangyu Sun^{1*}

Fan Long⁶



PipeZK: Accelerating Zero-Knowledge Proof with a Pipelined Architecture

Ye Zhang^{1,5} Shuo Wang¹ Xian Zhang³ Jiangbin Dong^{4,7} Xingzhong Mao⁷ Fan Long⁶
Cong Wang⁸ Dong Zhou² Mingyu Gao^{2,7*} Guangyu Sun^{1*}

PipeMSM: Hardware Acceleration for Multi-Scalar Multiplication

Charles. F. Xavier
charlie@ingonyama.com



PipeZK: Accelerating Zero-Knowledge Proof with a Pipelined Architecture

Ye Zhang^{1,5} Shuo Wang¹ Xian Zhang³ Jiangbin Dong^{4,7} Xingzhong Mao⁷ Fan Long⁶
Cong Wang⁸ Dong Zhou² Mingyu Gao^{2,7*} Guangyu Sun^{1*}

PipeMSM: Hardware Acceleration for Multi-Scalar Multiplication

Charles. F. Xavier
charlie@ingonyama.com

FPGA Acceleration of Multi-Scalar Multiplication: CycloneMSM

Kaveh Aasaraai, Don Beaver, Emanuele Cesena, Rahul Maganti,
Nicolas Stalder and Javier Varela



PipeZK: Accelerating Zero-Knowledge Proof with a Pipelined Architecture

Ye Zhang^{1,5} Shuo Wang¹ Xian Zhang³ Jiangbin Dong^{4,7} Xingzhong Mao⁷ Fan Long⁶
Cong Wang⁸ Dong Zhou² Mingyu Gao^{2,7*} Guangyu Sun^{1*}

PipeMSM: Hardware Acceleration for Multi-Scalar Multiplication

Charles. F. Xavier
charlie@ingonyama.com

FPGA Acceleration of Multi-Scalar Multiplication: CycloneMSM

Kaveh Aasaraai, Don Beaver, Emanuele Cesena, Rahul Maganti,
Nicolas Stalder and Javier Varela

cuZK: Accelerating Zero-Knowledge Proof with A Faster Parallel Multi-Scalar Multiplication Algorithm on GPUs

Tao Lu, Chengkun Wei, Ruijing Yu, Yi Chen, Li Wang, Chaochao Chen,
Zeke Wang, and Wenzhi Chen



ZPRIZE

Prizes

* General interest prize category for public goods which benefit multiple protocols/proof systems

OPEN DIVISION

Accelerating MSM Operations
on GPU/FPGA

OPEN DIVISION

Accelerating NTT Operations on
an FPGA

OPEN DIVISION

Plonk-DIZK GPU Acceleration

PRIZE
AMOUNT

\$1,250,000 USD ✓



Why Hardware Acceleration

A brief detour into gaming



Quake, 1997



Call of Duty, 2022



The difference - GPU



Multiplayer games



Multiplayer games

- Two ways to run a multiplayer game: (1) client-server, (2) P2P



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server
 - Ownership
 - Expensive
 - Scalability
 - Privacy



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server
 - Ownership
 - Expensive
 - Scalability
 - Privacy
- **P2P** - players can cheat reporting on their state



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server
 - Ownership
 - Expensive
 - Scalability
 - Privacy
- **P2P** - players can cheat reporting on their state

ZK fixes this 💪



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server
 - Ownership
 - Expensive
 - Scalability
 - Privacy
- **P2P** - players can cheat reporting on their state

ZK fixes this 💪



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server

- Ownership
- Expensive
- Scalability
- Privacy

ZK fixes this 💪

- **P2P** - players can cheat reporting on their state



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server

- Ownership
- Expensive
- Scalability
- Privacy

ZK fixes this 💪

- **P2P** - players can cheat reporting on their state

Problem: we need to run many ZKPs per sec for state changes



Multiplayer games



- Two ways to run a multiplayer game: (1) client-server, (2) P2P
- **Client-server (common)** - centralized (trusted) authoritative server

- Ownership
- Expensive
- Scalability
- Privacy

ZK fixes this 💪

- **P2P** - players can cheat reporting on their state

Happy to Discuss!

Problem: we need to run many ZKPs per sec for state changes



Are we there yet?



Are we there yet?

- **CryptoPunk ownership:** 25[s] proof generation time [1]
- **RISC-V ZKVM:** 30k - 1M instructions per second (70's computer speed)[2]
- **ZK MLaaS of ImageNet scale inference:** 2457[s] proving time at <80% accuracy [3]
- **Scroll ZKEVM:** 2535[s] for EVM circuit prover [4]
- **zkBridge** costs 8100 USD per month to operate [5]



Are we there yet?

- **CryptoPunk ownership:** 25[s] proof generation time [1]
- **RISC-V ZKVM:** 30k - 1M instructions per second (70's computer speed)[2]
- **ZK MLaaS of ImageNet scale inference:** 2457[s] proving time at <80% accuracy [3]
- **Scroll ZKEVM:** 2535[s] for EVM circuit prover [4]
- **zkBridge** costs 8100 USD per month to operate [5]

Answer: We are not far!



ZPU

Zero-Knowledge Processing Unit



ZPU

Zero-Knowledge Processing Unit

1. What goes into a ZPU? (hint: standardization can help)
2. What can be done until we have ZPUs?





ZPU


How to identify one when you see one?



ZPU

How to identify one when you see one?

- AES-NI but for ZK
 - Fixed standardized algorithm vs. moving, evolving target
- A dedicated CPU
 - Good for integer and bitwise arithmetic, we work with finite fields





ZPU

How to identify one when you see one?

- Chip for modular arithmetic {add, mul, inv}
 - What field? Can we add specific primitives?



ZPU

How to identify one when you see one?

- Chip for modular arithmetic {add, mul, inv}
 - What field? Can we add specific primitives?
- Polynomial arithmetic over finite field(s) (FFT acceleration)
- Elliptic curve(s) operations acceleration (MSM acceleration)
- Vectorized arithmetic (Hash function(s) acceleration)



ZPU

How to identify one when you see one?

- Chip for modular arithmetic {add, mul, inv}
 - What field? Can we add specific primitives?
- Polynomial arithmetic over finite field(s) (FFT acceleration)
- Elliptic curve(s) operations acceleration (MSM acceleration)
- Vectorized arithmetic (Hash function(s) acceleration)

Small area
Less performance
Less Power



Large area
More performance
More Power

Prioritized list for standardization for ZKProof6

Opinionated list!

- Finite Field/s arithmetic:
 - (e.g. 64bit prime vs. 256bit prime)
 - Elliptic Curve
- NTT-less (STARK vs. SNARK)
 - Standardize Protocol/s
- Vector operations:
 - ZK-friendly (arithmetic) hash function/s
 - Arithmetization (Plonk style vs. R1CS style)
- Range of circuit size



Prioritized list for standardization for ZKProof6

Opinionated list!

- Finite Field/s arithmetic:
 - (e.g. 64bit prime vs. 256bit prime)
 - Elliptic Curve
- NTT-less (STARK vs. SNARK)
 - Standardize Protocol/s
- Vector operations:
 - ZK-friendly (arithmetic) hash function/s
 - Arithmetization (Plonk style vs. R1CS style)
- Range of circuit size

Happy to Discuss!



What can be done today?



What can be done today?

- Do Science
- Build solutions on existing Hardware



Research Questions

A selection



Research Questions

A selection

- What is the optimal HW-Friendly ZK protocol?



Research Questions

A selection

- What is the optimal HW-Friendly ZK protocol?
- Can we design a hash function that is both ZK-friendly and HW-friendly?



Research Questions

A selection

- What is the optimal HW-Friendly ZK protocol?
- Can we design a hash function that is both ZK-friendly and HW-friendly?
- Imagine a world where prover computing is no longer a bottleneck... now what?



Research Questions

A selection

- More HW-Friendly ZK protocols
- ZK-friendly & HW-friendly hash functions
- Imagine a world where prover computing is no longer a bottleneck... now what?
- More Memory awareness/management in SNARKs/STARKSs



Research Questions

A selection

- More HW-Friendly ZK protocols
- ZK-friendly & HW-friendly hash functions
- Imagine a world where prover computing is no longer a bottleneck... now what?
- More Memory awareness/management in SNARKs/STARKSs

Happy to Discuss!



Build



- Available Hardware
- What to build?
- Examples

Available Hardware



GPUs

FPGAs



Available Hardware



GPUs

- Cheap to buy

FPGAs

Available Hardware



GPUs

- Cheap to buy

FPGAs

- Cheap to operate

Available Hardware



GPUs

- Cheap to buy
- Mining GPUs can be repurposed to run ZK

FPGAs

- Cheap to operate

Available Hardware



GPUs

- Cheap to buy
- Mining GPUs can be repurposed to run ZK

FPGAs

- Cheap to operate
- FPGA data centers can be utilized

Available Hardware



GPUs

- Cheap to buy
- Mining GPUs can be repurposed to run ZK
- Available on-demand in the cloud

FPGAs

- Cheap to operate
- FPGA data centers can be utilized
- Available on-demand in the cloud

Name	FPGAs	vCPUs	Instance Memory (GiB)	SSD Storage (GB)	Enhanced Networking	EBS Optimized	On-Demand Price/hr*
f1.2xlarge	1	8	122	470	Yes	Yes	\$1.65
f1.4xlarge	2	16	244	940	Yes	Yes	\$3.30
f1.16xlarge	8	64	976	4 x 940	Yes	Yes	\$13.20

Available Hardware



GPUs

- Cheap to buy
- Mining GPUs can be repurposed to run ZK
- Available on-demand in the cloud
- Empirically: Gaming GPUs perform better
 - Cons: Form factor, Energy

FPGAs

- Cheap to operate
- FPGA data centers can be utilized
- Available on-demand in the cloud

Name	FPGAs	vCPUs	Instance Memory (GiB)	SSD Storage (GB)	Enhanced Networking	EBS Optimized	On-Demand Price/hr*
f1.2xlarge	1	8	122	470	Yes	Yes	\$1.65
f1.4xlarge	2	16	244	940	Yes	Yes	\$3.30
f1.16xlarge	8	64	976	4 x 940	Yes	Yes	\$13.20

Available Hardware



GPUs

- Cheap to buy
- Mining GPUs can be repurposed to run ZK
- Available on-demand in the cloud
- Empirically: Gaming GPUs perform better
 - Cons: Form factor, Energy
- Open-source libraries: e.g. sppark, ec-gpu

ec-gpu & ec-gpu-gen

crates.io v0.2.0 docs passing circleci passing rustc 1.54+

crates.io v0.4.0 docs passing circleci passing rustc 1.54+

CUDA/OpenCL code generator for finite-field arithmetic over prime fields and elliptic curve arithmetic constructed with Rust.

FPGAs

- Cheap to operate
- FPGA data centers can be utilized
- Available on-demand in the cloud

Name	FPGAs	vCPUs	Instance Memory (GiB)	SSD Storage (GB)	Enhanced Networking	EBS Optimized	On-Demand Price/hr*
f1.2xlarge	1	8	122	470	Yes	Yes	\$1.65
f1.4xlarge	2	16	244	940	Yes	Yes	\$3.30
f1.16xlarge	8	64	976	4 x 940	Yes	Yes	\$13.20

- Open-source

cloud-ZK

A toolkit for developing ZKP acceleration in the cloud

Build



- **Infra & Foundations**
 - Prover-as-a-Service
 - Decentralized Prover-as-a-Service
 - ZK marketplace
 - App-specific, e.g. Filecoin miner, Danksharding Builder
- **Frameworks Integrations**, e.g. Arkworks, Circom

Example 1: Aleo miner



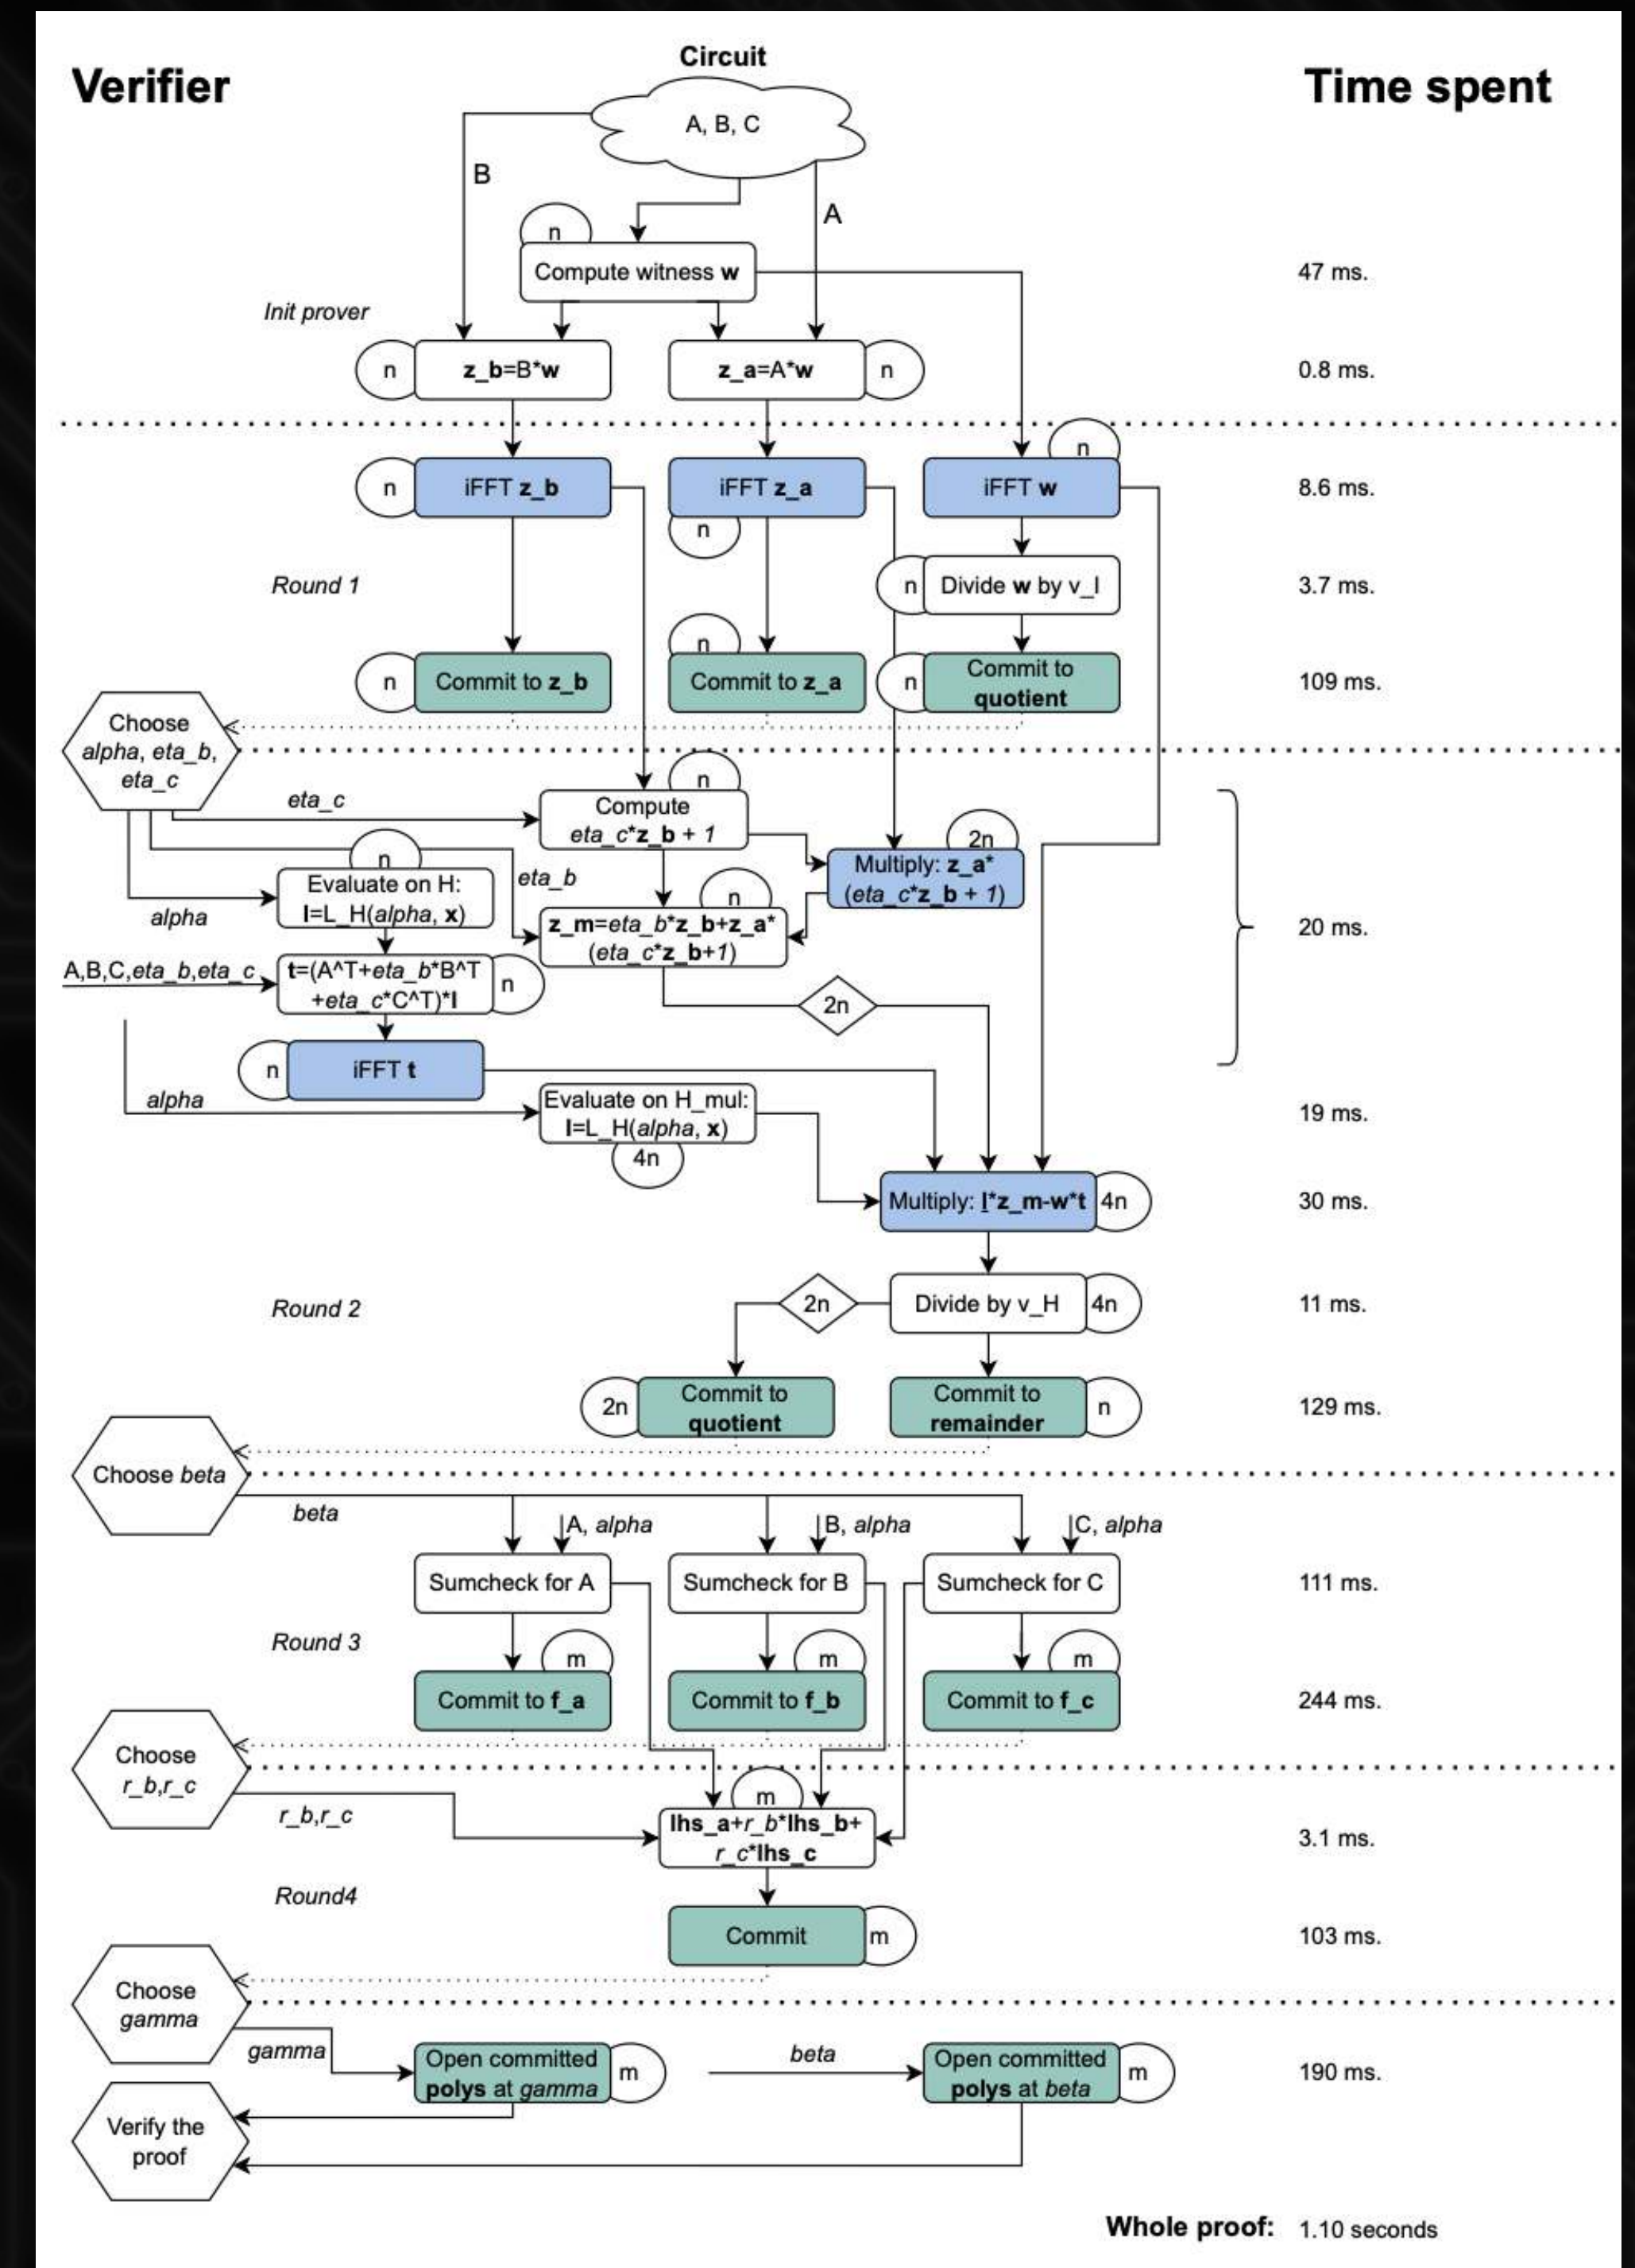
Example 1: Aleo miner

- GPU/FPGA/GPU+FPGA
- Electricity cost is important factor



Example 1: Aleo miner

- GPU/FPGA/GPU+FPGA
- Electricity cost is important factor
- Optimized Marlin Prover
 - BLS12-377
 - Mostly MSMs (70%) and FFTs (15%)



Example 1: Aleo miner

- GPU/FPGA/GPU+FPGA
 - Electricity cost is important factor
- Optimized Marlin Prover
 - BLS12-377
 - Mostly MSMs (70%) and FFTs (15%)

CPU: 0.9 [proofs/sec]
HW: 8 [proofs/sec]



Example 2: Rapid RapidSnark



Example 2: Rapid RapidSnark

rapidsnark


rapid snark is a zkSnark proof generation written in C++ and intel assembly. That generates proofs created in [circom](#) and [snarkjs](#) very fast.

<https://github.com/iden3/rapidsnark>



Example 2: Rapid RapidSnark

- A popular tool chain for ZK applications
- Optimized Groth 16
 - BN254
 - 4 MSMs with G1, 1 MSM with G2, rest are NTTs
- Dev-ops: Beefy CPU instance.
- Rapid RapidSnark
 - replace CPU with GPU/FPGA in the cloud (AWS)
 - same cost (2 USD/h) or less



zkAttestor: Benchmarks

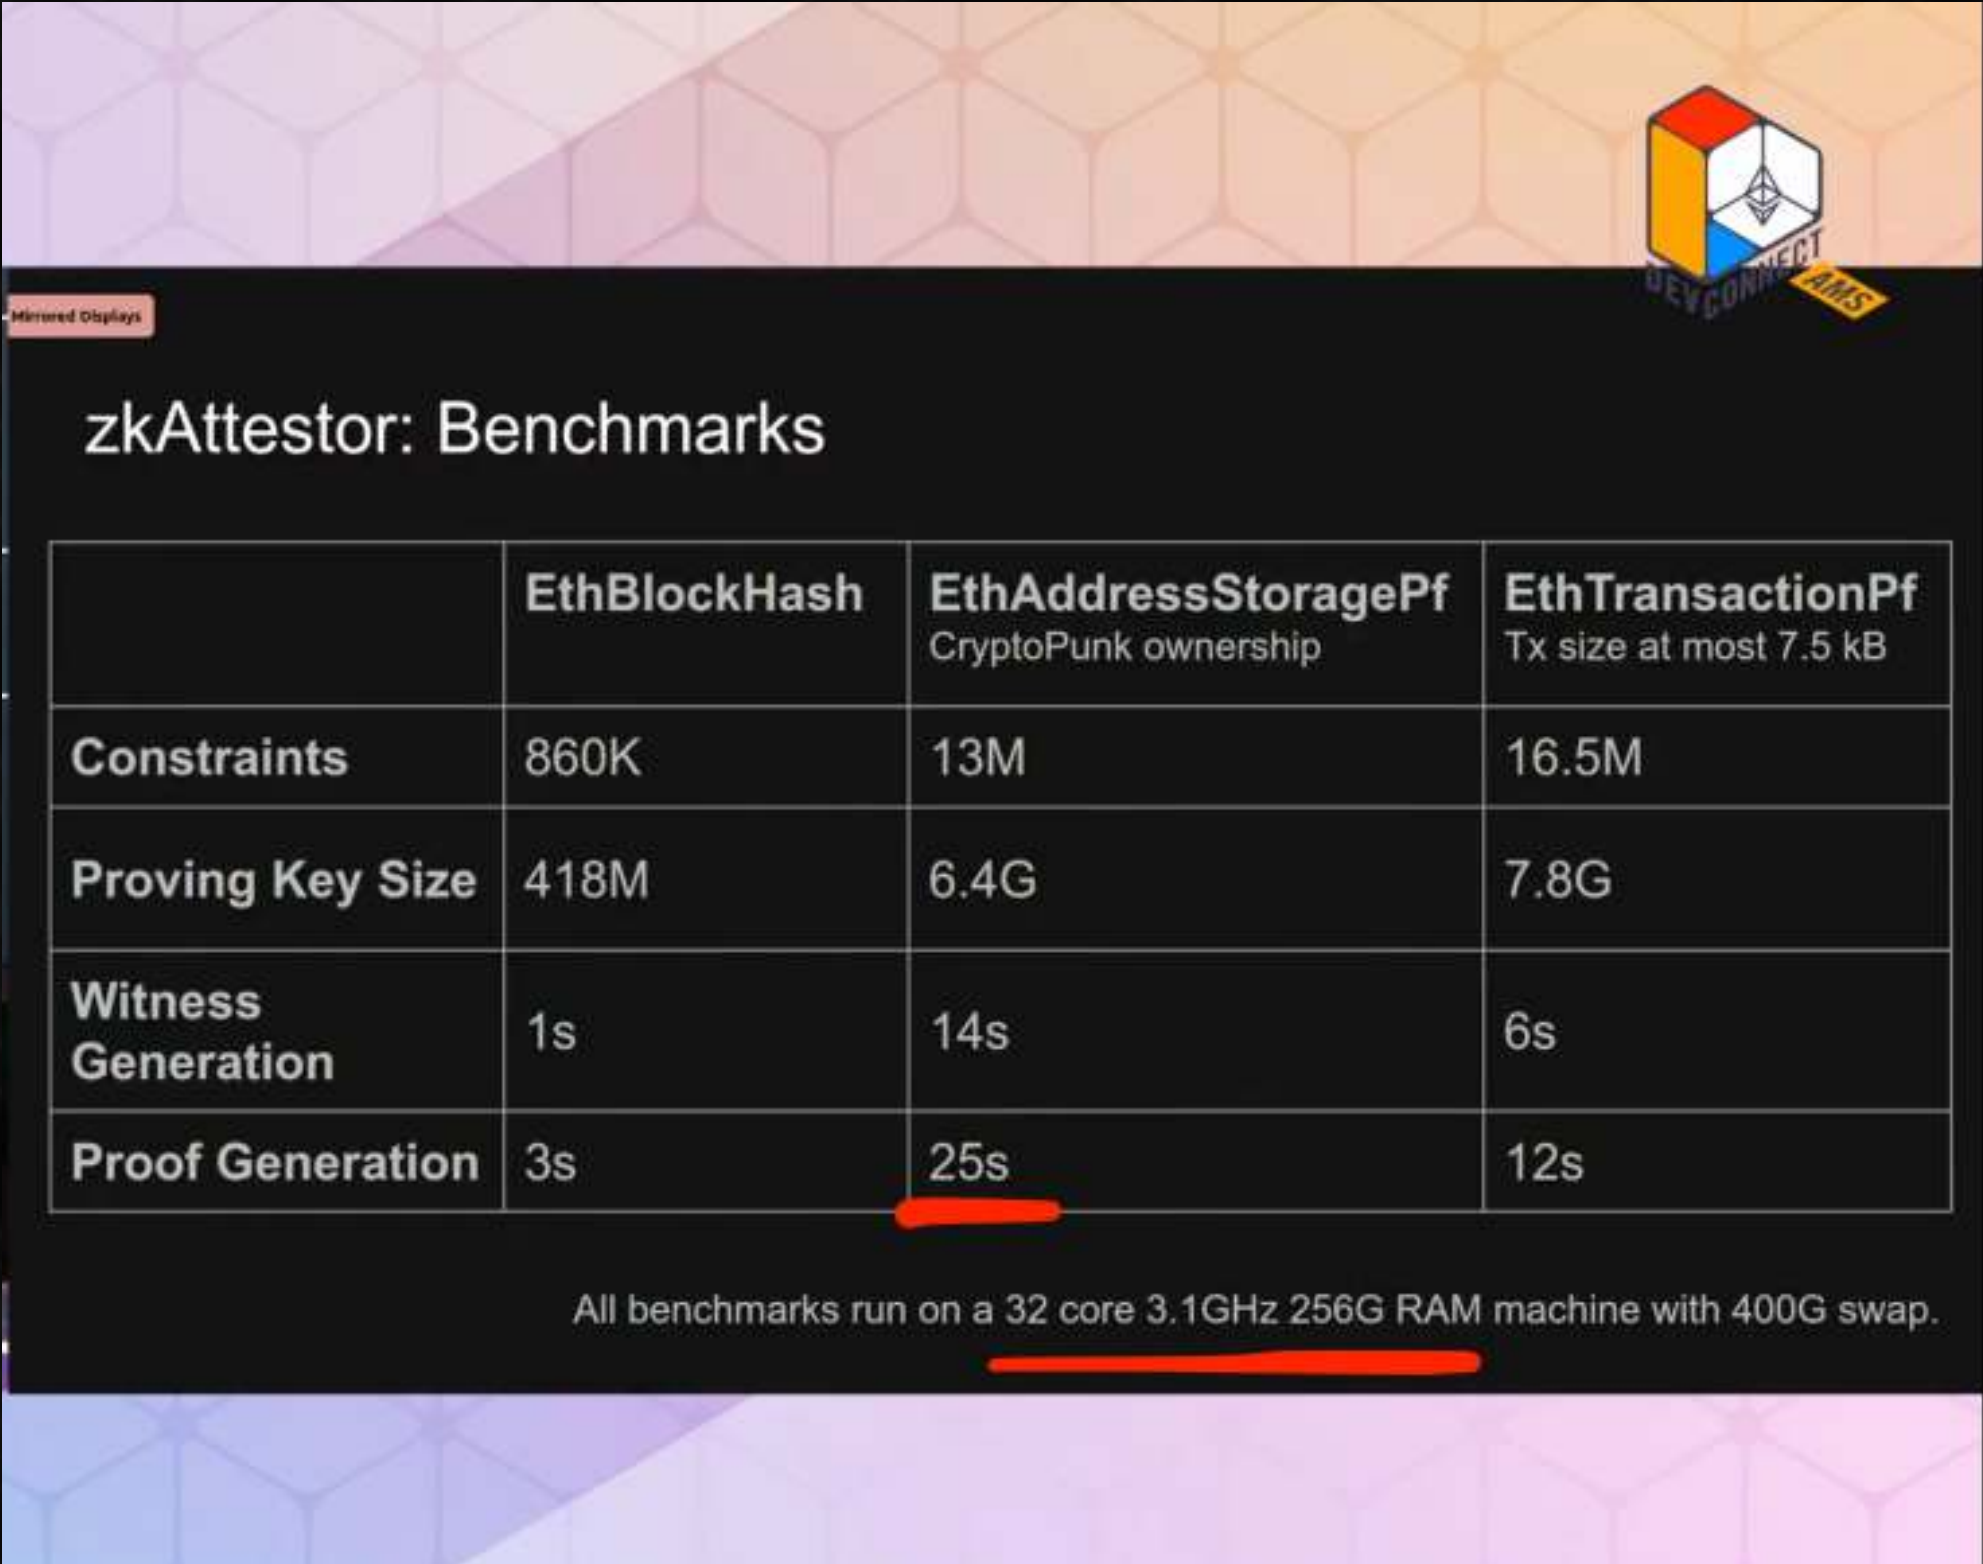
	EthBlockHash	EthAddressStoragePf CryptoPunk ownership	EthTransactionPf Tx size at most 7.5 kB
Constraints	860K	13M	16.5M
Proving Key Size	418M	6.4G	7.8G
Witness Generation	1s	14s	6s
Proof Generation	3s	25s	12s

All benchmarks run on a 32 core 3.1GHz 256G RAM machine with 400G swap.



Example 2: Rapid RapidSnark

- A popular tool chain for ZK applications
- Optimized Groth 16
 - BN254
 - 4 MSMs with G1, 1 MSM with G2, rest are NTTs
- Dev-ops: Beefy CPU instance.
- Rapid RapidSnark
 - replace CPU with GPU/FPGA in the cloud (AWS)
 - same cost (2 USD/h) or less



The image shows a screenshot of a benchmarking tool interface. At the top right is a logo for 'DEVCONNECT' with a cube icon. Below it, the title 'zkAttestor: Benchmarks' is displayed. A table follows with four columns: an empty header cell, 'EthBlockHash', 'EthAddressStoragePf' (with 'CryptoPunk ownership' as a subtitle), and 'EthTransactionPf' (with 'Tx size at most 7.5 kB' as a subtitle). The rows are 'Constraints', 'Proving Key Size', 'Witness Generation', and 'Proof Generation'. The 'Proof Generation' row has a red underline under the '25s' value. At the bottom, a note states: 'All benchmarks run on a 32 core 3.1GHz 256G RAM machine with 400G swap.' with a red underline.

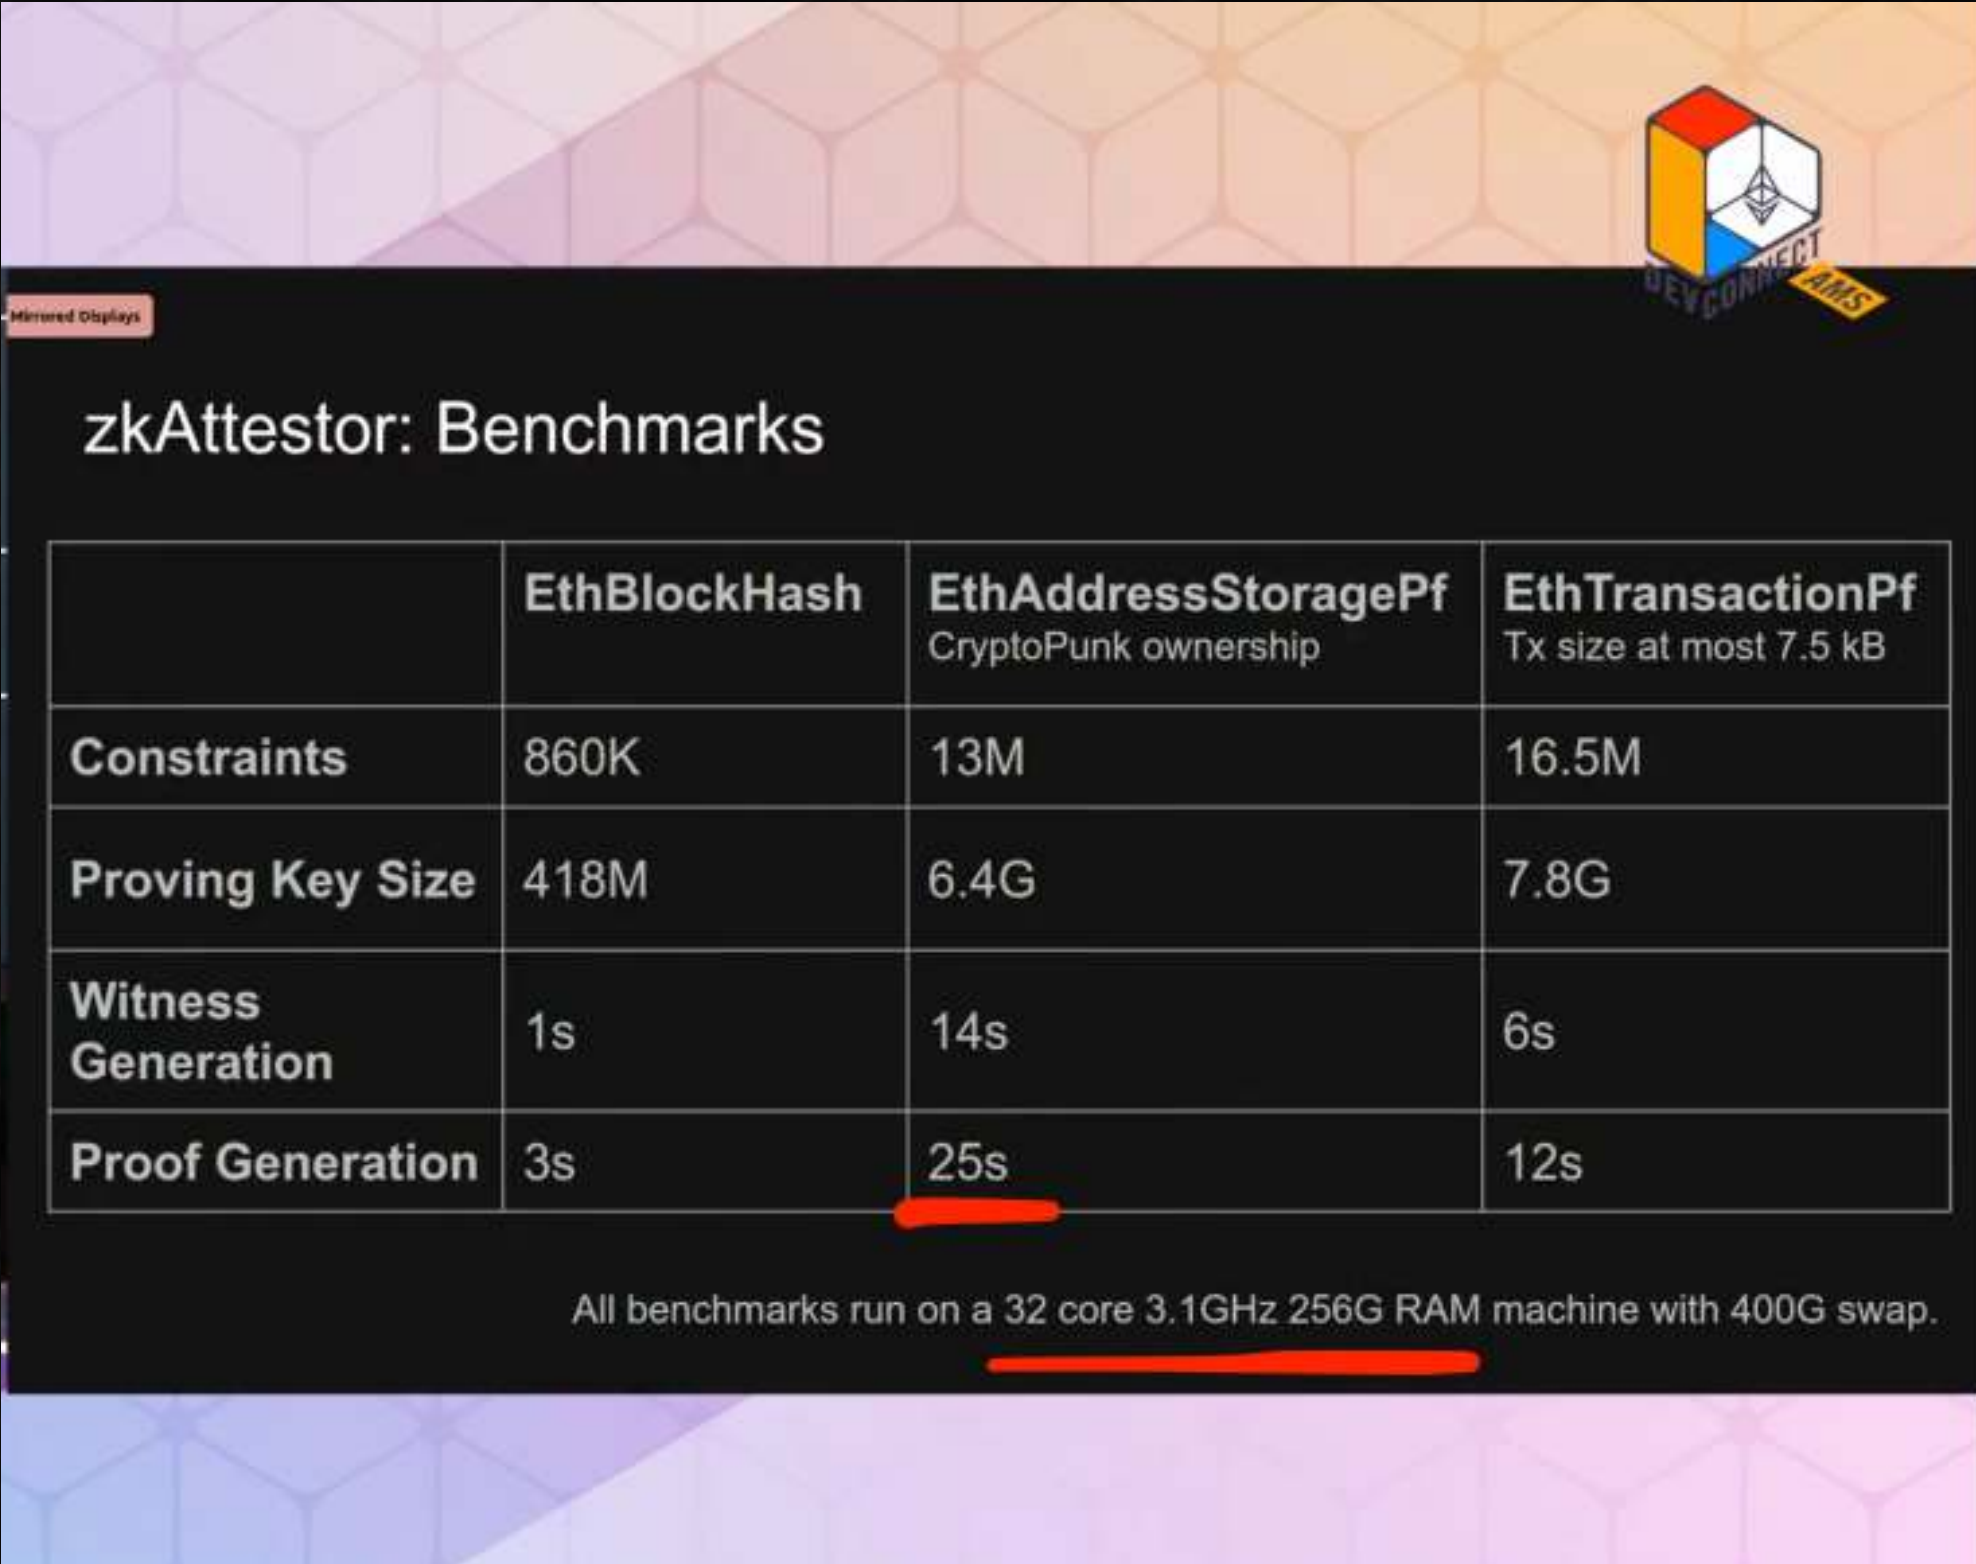
	EthBlockHash	EthAddressStoragePf CryptoPunk ownership	EthTransactionPf Tx size at most 7.5 kB
Constraints	860K	13M	16.5M
Proving Key Size	418M	6.4G	7.8G
Witness Generation	1s	14s	6s
Proof Generation	3s	25s	12s

All benchmarks run on a 32 core 3.1GHz 256G RAM machine with 400G swap.



Example 2: Rapid RapidSnark

- Accelerated: G_1, G_2, NTT
- Challenges:
 - Two elliptic curves
 - Hardware is not fully utilized
 - NTT grows $O(n \log(n))$, MSM grows $O(n)$
- Metric: Latency, Throughput
- Result: 5x improvement (open source soon)



The image shows a screenshot of a benchmarking website for zkAttestor. The page has a dark theme with a light blue and orange header. A logo for 'DEVCONNECT' is in the top right. The title 'zkAttestor: Benchmarks' is centered. Below it is a table with four columns: 'EthBlockHash', 'EthAddressStoragePf' (with subtext 'CryptoPunk ownership'), and 'EthTransactionPf' (with subtext 'Tx size at most 7.5 kB'). The rows are 'Constraints', 'Proving Key Size', 'Witness Generation', and 'Proof Generation'. The 'Proof Generation' row has a red underline under the '25s' value. At the bottom, a note states: 'All benchmarks run on a 32 core 3.1GHz 256G RAM machine with 400G swap.' with a red underline.

	EthBlockHash	EthAddressStoragePf CryptoPunk ownership	EthTransactionPf Tx size at most 7.5 kB
Constraints	860K	13M	16.5M
Proving Key Size	418M	6.4G	7.8G
Witness Generation	1s	14s	6s
Proof Generation	3s	25s	12s

All benchmarks run on a 32 core 3.1GHz 256G RAM machine with 400G swap.



Thank you



INGONYAMA

Omer Shlomovits , ZKProof 5 Workshop